
Force sketching in MD: Can we replace multi-head committees with screening methods for faster force uncertainty compute?

Anonymous Author(s)

Affiliation

Address

email

Abstract

Machine-learned interatomic potentials simulate molecules at a small fraction of the cost of quantum chemistry, but they are only useful if they can tell us when they are guessing. The standard way to ask is to train a *committee* of predictors and measure how much they disagree. A *multi-head* committee makes this almost free for energies: M prediction heads share one network body, so a single forward pass through the network returns all M energy estimates at once. Forces are not free. A force is the derivative of the energy with respect to the atomic positions, and recovering each head’s forces takes its own backward pass through that network — so force disagreement, the quantity that actually decides which structures a simulation should stop and verify, costs roughly M times more than energy disagreement.

We ask whether a handful of cheap random probes can estimate that disagreement without computing all M sets of forces, using a pretrained eight-head committee on four molecular systems. The answer is no for *replacement* and yes for *screening*. Even given six of the seven probes the exact computation would use, the best estimator identifies only 72–80% of the genuinely most uncertain structures, so an approximate score cannot stand in for the exact one. But a calibrated filter built from that same score safely skips 84–87% of the expensive computations while still catching 96–98% of the uncertain structures, and beats cheaper alternatives in 11 of 12 comparisons. Whether this saves wall-clock time turns out to depend on how the model is run: filtering is worth $1.85\times$ when structures are processed in batches, but only $1.11\times$ one at a time, where it becomes a net slowdown. The decision-quality result is robust; the speed result is not. We report this as a failure mode rather than a fix: the estimator that looks correct in expectation is the one that fails at the decision, and it is calibration, not a better estimator, that recovers a usable guarantee.

1 Introduction

Uncertainty quantification is what makes learned interatomic potentials usable beyond their training distribution: it drives active learning, flags extrapolation during molecular dynamics, and decides when a structure must go back to the quantum-chemistry calculation the model was trained to imitate, orders of magnitude slower and the very thing we are avoiding. Committee disagreement remains the most widely used estimator of that error: where training data pinned the model down its members agree, and where it did not they diverge, so their spread stands in for an error we cannot see [Behler, 2015, Podryabinkin and Shapeev, 2017, Carrete et al., 2023]. Multi-head committees [Beck et al., 2025] — shallow ensembles more broadly [Kellner and Ceriotti, 2024] — make that cheap by

attaching M readout heads to one shared message-passing trunk, so all M energies come from a single forward pass instead of M separately trained models.

What this paper does and does not show. The failure we report is a failure of *unbiasedness as a proxy for usefulness*. Every estimator below is unbiased for the force-variance statistic, with closed-form finite- K corrections, and each is still unusable as a drop-in replacement, because the decision downstream is a rank or a threshold rather than a mean and an unbiased score with the wrong noise structure gets that decision wrong. Two limits are worth stating plainly. The systems are small drug-like molecules of 9–24 atoms, not proteins or assemblies, so the extreme-value bias we characterise over $3N$ noisy per-component estimates should be read as a lower bound on what the same estimator would suffer at biomolecular scale, where N is orders of magnitude larger. And the one clearly positive result, the calibrated gate, inherits the coverage guarantee of split conformal prediction, which holds under exchangeability; a molecular dynamics trajectory is not exchangeable, so the guarantee applies to the held-out frames we evaluate on and not automatically to frames drawn sequentially from a running simulation.

That amortization does not extend to forces. As the multi-head committee work states directly, “a single backward pass can not yield the forces for all the heads ... the standard deviation requires multiple reverse passes through the computational graph.” Forces are what MD and most acquisition rules actually consume, so the expensive half of committee uncertainty is exactly the half the shared trunk fails to amortize.

Downstream, though, one rarely needs every member force: acquisition needs a *ranking*, a safety gate needs a *threshold decision*, error analysis needs the most uncertain *atom*. That motivates estimating the force-variance statistic directly rather than reconstructing the member forces defining it. We call this ForceSketch: replace the exact $(M-1)$ -direction centered head basis with K randomized or orthogonal directions, each still costing one reverse pass.

Contributions. We formulate multi-head force disagreement as a statistic of the centered energy Jacobian and show the exact computation needs only the centered $(M-1)$ -dimensional head subspace; we give unbiased Gaussian, Haar-orthogonal, Rademacher and pairwise estimators with closed-form finite- K corrections, correcting a removable singularity in the Haar constant. We then report a *negative* result we believe is the useful one — at $K \leq 4$ none of these preserves the exact top-5% acquisition set, on any of four systems — and show what does work: an empirically calibrated screening gate over a low-rank control variate, whose decision quality and wall-clock saving we quantify separately, because only the first survives.

2 Problem formulation

Setup and notation. A structure of N atoms has Cartesian coordinates $x \in \mathbb{R}^D$ with $D = 3N$. We index a single coordinate by $d \in \{1, \dots, D\}$, an atom by $a \in \{1, \dots, N\}$, and the three Cartesian components of one atom by $\gamma \in \{1, 2, 3\}$, so the pair (a, γ) names the same thing as d . The committee has M heads: M small readout networks attached to one shared message-passing trunk. Because they share the trunk, a single forward pass returns all M head energies at once, as the vector $E(x) = [E_1(x), \dots, E_M(x)]^\top \in \mathbb{R}^M$. Head m ’s force field is the negative energy gradient $f_m = -\nabla_x E_m \in \mathbb{R}^D$, and we collect the M of them as the columns of

$$F(x) = [f_1 \cdots f_M] \in \mathbb{R}^{D \times M}, \quad \text{so } F_{dm} \text{ is head } m\text{’s force along coordinate } d.$$

Write $\bar{F}_d = \frac{1}{M} \sum_{m=1}^M F_{dm}$ for the committee-mean force on coordinate d — the force the simulation actually uses.

Two disagreement scores, at two scales. Uncertainty is how much the heads disagree. On a *single coordinate* that is their sample variance; summing over coordinates gives one number for the *whole structure*:

$$\underbrace{v_d}_{\text{per coordinate}} = \frac{1}{M-1} \sum_{m=1}^M (F_{dm} - \bar{F}_d)^2, \quad \underbrace{S}_{\text{per structure}} = \sum_{d=1}^D v_d = \frac{\|FP\|_F^2}{M-1}. \quad (1)$$

Here $\mathbf{1} \in \mathbb{R}^M$ is the all-ones vector, $P = I_M - \frac{1}{M} \mathbf{1}\mathbf{1}^\top$ is the projector that subtracts the head mean (*centering*), and $\|\cdot\|_F$ is the Frobenius norm. Centering is what turns raw predictions into




Figure 1: What ForceSketch replaces. All M head energies come from one forward pass, but forces do not: exact disagreement needs the full centered head subspace, M lanes. ForceSketch spends $K+1$, and — since that is not accurate enough to *replace* exact uncertainty at $K \leq 4$ under either acquisition rule (Section 5.1) — its score feeds a calibrated gate that runs the exact path only where the bound cannot rule out the threshold. Whether fewer lanes also cost less time depends on the autodiff regime (Section 5.2).

disagreement: FP holds each head’s deviation from the consensus, and v_d is just the squared length of its d -th row, $v_d = \|(FP)_d\|_2^2 / (M - 1)$. Pooling the three components of one atom gives the two atom-level scores used in the literature, $u_a^{\text{RMS}} = (\frac{1}{3} \sum_{\gamma} v_{a\gamma})^{1/2}$ and $u_a^{\text{MHC}} = \frac{1}{3} \sum_{\gamma} \sqrt{v_{a\gamma}}$, where $v_{a\gamma}$ is v_d for the coordinate $d = (a, \gamma)$. Section 4 says which of these the acquisition rule ranks by.

Why energies are cheap and forces are not. The M energies cost one forward pass because the heads share a trunk. Forces are gradients, and reverse-mode differentiation returns the gradient of *one scalar* per backward pass. Feeding it the mean energy yields the mean force, not the spread; feeding it each E_m in turn yields all M force fields but costs M backward passes. We call one backward pass a *lane*.

The saving grace is that S depends on F only through the centered matrix FP , whose columns sum to zero and therefore span at most $M - 1$ dimensions. Let $Q \in \mathbb{R}^{M \times (M-1)}$ have orthonormal columns $q_1, \dots, q_{M-1}$ spanning exactly that centered subspace:

$$Q^T Q = I_{M-1}, \quad Q^T \mathbf{1} = 0, \quad Q Q^T = P.$$

Each column is one lane’s worth of work. Differentiating the scalar $q_j^T E$ — a fixed weighted combination of the head energies — costs a single backward pass and returns

$$g_j = -\nabla_x (q_j^T E) = F q_j \in \mathbb{R}^D, \quad \text{and} \quad v_d = \frac{1}{M-1} \sum_{j=1}^{M-1} g_{j,d}^2. \quad (2)$$

So the exact variance needs $M - 1$ lanes, plus one more along the uniform direction $\frac{1}{M} \mathbf{1}$ for the mean force the simulation needs anyway: M lanes in total, matching the naive per-head loop. *Our question is whether downstream uncertainty decisions really need all $M - 1$ centered lanes.*

3 Method

Gaussian and Haar sketches. Our estimators are the molecular-UQ instance of randomized trace and norm estimation [Hutchinson, 1990, Avron and Toledo, 2011, Meyer et al., 2021]; the orthogonal variant is the head-space analogue of orthogonal random features [Yu et al., 2016], and the control

104 variate plays the role that a low-rank sketch plays in randomized linear algebra [Halko et al., 2011].
 105 Unbiased approximate vector–Jacobian products have also been used to cheapen *training* [Bakong
 106 et al., 2026], by sparsifying the backward pass inside the network; we instead compress the *output-*
 107 *head query space* at inference, leaving each reverse pass exact and reducing how many are needed to
 108 estimate an existing committee statistic. The idea is to stop asking for all $M - 1$ centered directions
 109 and instead *probe* the centered subspace with $K \ll M - 1$ random ones, then rescale. Each probe
 110 still costs exactly one lane, so the question is how much of the disagreement K probes recover.

111 Concretely, draw $z_k \sim \mathcal{N}(0, I_M)$ and center it, $w_k = Pz_k$; one lane returns $g_k = Fw_k \in \mathbb{R}^D$, and

$$\hat{v}_d = \frac{1}{K(M-1)} \sum_{k=1}^K g_{k,d}^2 \quad (3)$$

112 is an unbiased estimator of v_d — on average it gives the right answer, though any single draw does
 113 not. The intuition is the standard one behind randomized trace estimation: a random probe sees each
 114 direction of the subspace equally in expectation, so the average of K probes estimates the total energy
 115 in all $M - 1$ of them.

116 Random probes can overlap, wasting budget. Sampling them *orthogonally* instead removes that waste.
 117 Let $O \in \mathbb{R}^{(M-1) \times K}$ have orthonormal columns drawn Haar-uniformly (that is, uniformly over all
 118 K -dimensional orientations, obtained in practice by orthonormalizing a Gaussian matrix), and set
 119 $W = QO \in \mathbb{R}^{M \times K}$: a random orthonormal K -frame lying inside the centered subspace. Collecting
 120 the K lanes it costs into $G = FW \in \mathbb{R}^{D \times K}$ — so G_{dk} is the response of coordinate d to probe k —
 121 gives

$$\hat{v}_d^{\text{Haar}} = \frac{1}{K} \sum_{k=1}^K G_{dk}^2, \quad (4)$$

122 which is unbiased and becomes *exact* at $K = M - 1$, where the frame is a full basis of the centered
 123 subspace. Writing $r = M - 1$ for its dimension, the two estimators have relative standard deviations
 124 $\sqrt{2/K}$ and $\sqrt{2(r-K)/(K(r+2))}$, so orthogonalization multiplies the estimator variance by
 125 exactly $(r-K)/(r+2)$ — a factor of 4/9 at $K=3, M=8$. It costs nothing extra, and it is the single
 126 largest methodological lever we find.

127 **Finite- K corrections.** Reporting a standard deviation means taking a square root, and the square
 128 root of an unbiased variance estimate is *not* an unbiased standard-deviation estimate: $\sqrt{\cdot}$ is concave,
 129 so $\sqrt{\hat{v}_d}$ is biased low, the more so the fewer probes we spend. The fix is a multiplicative constant
 130 depending only on K (and, for orthogonal probes, on r). For Gaussian probes $K\hat{v}_d/v_d \sim \chi_K^2$ and the
 131 correction is $c_K = \sqrt{2/K} \Gamma(\frac{K+1}{2})/\Gamma(\frac{K}{2})$. For Haar probes the analogous constant is conventionally
 132 written with Beta functions, which is singular at $K = r$ (it evaluates $B(K/2, 0)$). Cancelling the
 133 shared $\Gamma(\frac{r-K}{2})$ gives a form that is finite everywhere and exactly 1 at full rank:

$$c_{K,r}^{\text{Haar}} = \sqrt{\frac{r}{K}} \frac{\Gamma(\frac{K+1}{2}) \Gamma(\frac{r}{2})}{\Gamma(\frac{K}{2}) \Gamma(\frac{r+1}{2})}, \quad c_{K,r}^{\text{Haar}} \rightarrow c_K \text{ as } r \rightarrow \infty. \quad (5)$$

134 We verify unbiasedness of $\sqrt{\hat{v}}/c$ to within Monte-Carlo error at 4×10^5 draws. These closed forms
 135 cover the Gaussian and Haar estimators; the control variate below mixes exact and sketched directions
 136 and has no such constant, so we apply the residual frame’s Haar correction, which is an approximation.
 137 It does not affect the screening results, because c_α is itself a multiplicative calibration and absorbs
 138 any constant scale error exactly.

139 **Low-rank control variate.** Split $P = Q_{r_0} Q_{r_0}^\top + R$, where $Q_{r_0} \in \mathbb{R}^{M \times r_0}$ has orthonormal columns
 140 inside the centered subspace and R is the projector onto what is left of it, the *residual* subspace of
 141 dimension $r - r_0$. The columns of Q_{r_0} are the r_0 leading eigenvectors of the *design-split* average
 142 $\bar{A} = \frac{1}{n} \sum_i P F_i^\top F_i P$ — an $M \times M$ matrix in *head* space, whose top eigenvector is the weighting
 143 of heads along which this committee disagrees most consistently. It is fitted separately for each
 144 system on structures disjoint from both the calibration and the test splits (Section 4). Fitting it on
 145 the calibration split instead would invalidate the conformal guarantee below, since split conformal
 146 requires the score function — learned basis included — to be frozen before the calibration ratios are
 147 formed. The idea is to spend a few lanes where the disagreement actually concentrates, and probe

only what is left over. Evaluate those r_0 leading directions exactly, costing $g_j = Fq_j$ for the columns q_j of Q_{r_0} , and sketch only the residual with $K' = K - r_0$ Haar probes drawn inside R 's range, collected as $G \in \mathbb{R}^{D \times K'}$ exactly as before but with the frame now living in the residual subspace:

$$\hat{v}_d = \frac{1}{r} \left[\sum_{j \leq r_0} g_{j,d}^2 + \frac{r-r_0}{K'} \sum_{k=1}^{K'} G_{dk}^2 \right], \quad (6)$$

which remains unbiased. By *spectrum* we mean the eigenvalues of $\bar{A}$ — how the total disagreement is divided among head-space directions. The motivation is then the *opposite* of the intuitive one: a concentrated spectrum does not explain why small K works, it is what *breaks* a random sketch, since a probe either catches the dominant direction or misses it.

Screening rather than replacement. If the sketch cannot *replace* the exact score, it can still say when the exact score is certainly small. Conformal prediction [Vovk et al., 2005, Angelopoulos et al., 2024] turns any score, however crude, into a calibrated bound using held-out data alone. On a calibration split of n structures disjoint from design and test, form the ratios $\eta_i = S(x_i)/(\hat{S}(x_i) + \epsilon)$, the factor by which the cheap score underestimates the exact one — and let c_α be their $\lceil (n+1)(1-\alpha) \rceil$ -th order statistic, large enough to cover all but an α fraction of them. Then $U(x) = c_\alpha \hat{S}(x)$ sits above $S(x)$ at least $1 - \alpha$ of the time, so $U(x) < \tau$ *certifies* that the exact score is below the high-uncertainty threshold τ too: skip when $U(x) < \tau$, fall back otherwise. Formally this is *marginal* coverage, $\mathbb{P}[S(x) \leq U(x)] \geq 1 - \alpha$, under exchangeability — not a conditional guarantee and not an out-of-distribution one. 3BPA at 1200 K calibrated against 300 K behaviour is exactly such a shift, so we report the observed coverage there rather than assume it.

4 Experimental setup

We use the pretrained eight-head MACE [Batatia et al., 2022] committees released with Beck et al. [2025]: two message-passing layers, 32 channels, $\ell_{\max} = 1$, $r_{\text{cut}} = 6 \text{ \AA}$. We evaluate the three head-distribution variants (*disjoint*, *overlapping*, *same*) on 3BPA [Kovács et al., 2021] at 1200 K (2139 structures, extrapolative relative to the 300 K training set) and the joint ten-molecule rMD17 [Christensen and von Lilienfeld, 2020] committee on ethanol (9 atoms), aspirin (21) and azobenzene (24), 1000 held-out structures each. Estimator evaluation is in float64 and this is not fastidiousness: the mean force exceeds the disagreement by a factor $\kappa \approx 36$ here, so a float32 “exact” reference carries more relative error than the 10^{-5} tolerance at which we verify full-rank exactness. Deployment timing is float32, on one RTX 5070 Laptop GPU, with PyTorch [Paszke et al., 2019] CUDA events measured per iteration; the fitted lane cost is $T(L) = 10.38 + 11.78 L$ ms. Ten random seeds are frozen before testing and none discarded; all intervals are paired bootstraps [Efron and Tibshirani, 1993] over structures with 10,000 resamples.

The acquisition decision, stated exactly. Every number below concerns one decision: from a pool of held-out structures, select the most force-uncertain fraction for exact reference computation. Following Beck et al. [2025], structures are ranked by the maximum disagreement of their force components; we evaluate that extreme-value rank in both of its defensible forms — maximum over components, maximum over atoms — and report the global sum as an easier diagnostic statistic:

$$\hat{u}^{\text{max-comp}}(x) = \max_d \hat{\sigma}_d, \quad \hat{u}^{\text{max-atom}}(x) = \max_a \hat{u}_a^{\text{MHC}}, \quad \hat{S}(x) = \sum_d \hat{v}_d. \quad (7)$$

A hat always denotes the sketched estimate of the corresponding exact quantity: the exact scores are the same three expressions with v_d and $\sigma_d = \sqrt{v_d}$ in place of $\hat{v}_d$ and $\hat{\sigma}_d$, so the exact S here is exactly the S of Eq. (1). The correction c is c_K for Gaussian probes and $c_{K,r}^{\text{Haar}}$ for orthogonal ones (Section 3); u_a^{MHC} is as in Section 2. We take $u^{\text{max-comp}}$ as primary, being the most literal reading, report $u^{\text{max-atom}}$ alongside it in the artifact, and treat S as the statistically *easier* secondary: a sum over $3N$ coordinates concentrates, whereas both maxima are extreme-value statistics of a noisy $\hat{\sigma}$.

The two experiments select positives differently and we do not conflate them. Section 5.1 measures **top-5% recall**: the 5% of evaluated structures with the largest *exact* score, ties resolved inclusively, so the set may slightly exceed 5%. Section 5.3 applies a *fixed* threshold instead, as a deployed gate must: τ is the 95th percentile of the exact score on the design split, and a test structure counts as

high-uncertainty when its exact score is at least τ — so the realised test prevalence is 0.028–0.050, not exactly 0.05. Every metric is computed *per seed and then averaged* over the ten frozen seeds: the deployed budget is K lanes, not $10K$, and averaging the sketches before scoring would be a different, far more accurate estimator than the one we claim.

Splits. Each system is partitioned once, under a frozen permutation, into three disjoint splits: a *design* split (20%) fitting the control-variate basis Q_{r_0} and the threshold τ ; a *calibration* split (20%) fitting the conformal constant c_α alone; and a *test* split (60%), touched once — 427/428/1284 structures for 3BPA and 200/200/600 per rMD17 molecule. The separation is not cosmetic: split-conformal coverage requires the score function to be frozen independently of the calibration observations, so a basis fitted on calibration structures would void the guarantee. The code asserts all three pairwise disjointness conditions and that the splits partition the set.

5 Results

5.1 Can force disagreement be sketched?

Partially, and not well enough to replace exact computation — under either acquisition rule, and markedly less well under the primary one.

Under $u^{\max\text{-comp}}$, Haar is the best sketch at every budget, but its top-5% recall is only 0.362–0.489 at $K=3$ and 0.484–0.570 at $K=4$ across the 4 systems. Rank correlation is far healthier than tail selection (Spearman 0.648–0.781 at $K=4$), which is the recurring pattern: a sketch that orders structures tolerably still fails to identify *which* ones land in the extreme tail. On the easier global statistic the same estimator reaches 0.529 [0.490, 0.566] and 0.619 [0.580, 0.648] at $K=3, 4$ on 3BPA — roughly ten points higher, and still far from substitution quality; Appendix B plots that easier view against K alongside the latency curve. Intervals are 95% paired bootstraps over structures (10,000 resamples) on the seed-averaged statistic; seed-to-seed standard deviations are 0.02–0.09 and are recorded separately.

Why the extreme-value rule is harder. Making each $\hat{\sigma}_d$ unbiased *marginally* is not the same as making $\max_d \hat{\sigma}_d$ unbiased: a maximum over $3N$ noisy estimates is biased upward, and the bias grows with the estimator’s noise. Measured directly, $u^{\max\text{-comp}}$ is inflated by 1.42–1.53 at $K=1$, falling to 1.07–1.08 at $K=4$ and to 1.000 at the full basis, where the estimator is exact and the bias must vanish — a useful check on the pipeline. This obstructs *replacement* but not screening (Section 5.3): the bias is monotone in the estimator’s noise, so a scale-calibrated gate absorbs it into c_α .

Orthogonalizing the probes is the largest single lever in the method. The estimator comparisons that follow are on the global statistic, where the paired-bootstrap machinery of Section 4 was frozen; the ordering they establish is a property of the estimators, not of the acquisition rule. The paired difference between Haar and Gaussian at $K=3$, evaluated on identical structure resamples, is 0.225 [0.186, 0.263] in top-5% recall, and is significant on all four systems. The low-rank control variate adds a further 0.111 [0.080, 0.145] over Haar at $K=4$, reaching 0.730 [0.692, 0.760] — still short of 0.90.

The head-space spectrum explains why. Writing $A = Q^\top \bar{F}^\top FQ$ for the averaged centered force-disagreement Gram matrix, the stable rank of the matrix we actually sketch is $\text{srank}(FQ) = \|FQ\|_F^2 / \|FQ\|_2^2 = \text{tr}(A) / \lambda_{\max} = 2.91\text{--}3.08$ of a maximum 7 across the four systems, with the leading direction carrying 0.325–0.344 of the disagreement (an isotropic spectrum would give 0.143). The spectrum is therefore *moderately diffuse rather than strongly rank-concentrated* — but far from isotropic, at 2.3–2.4 $\times$ the isotropic leading share. That is exactly the regime in which a random K -dimensional projection pays a variance penalty, and it is why projecting out the leading directions (the control variate) recovers so much.

Does sketching beat simply using fewer heads? It depends on whether an exact mean force is required, and the two framings give opposite answers, so Appendix F reports both rather than settling the comparison by definition.




Figure 2: Both panels use the *primary* max-component rule. Left, the negative result: top-5% recall against the sketch budget K for the best estimator, on all four systems. Recall rises steadily but does not reach the 0.90 replacement line at any budget short of the exact basis at $K=7$; error bars are the standard deviation across the ten frozen seeds. Right, the positive one: the screening decision for all four gates at matched lane budget. Up-and-right is better on both axes, so Pareto dominance is read directly. The control variate (red) sits above and to the right of every other gate on every system but one — on ethanol the free energy gate reaches recall 1.000 while skipping barely a third of evaluations. Bars are 95% paired bootstrap intervals; recall is the noisier axis because the positive sets hold only tens of structures.

242 5.2 Does reducing directions reduce GPU time?

243 Only in some regimes, and the honest answer required first making the strongest baseline available.
 244 Batched reverse mode — vectorizing the backward over cotangents while reusing one forward graph
 245 — is the baseline fairness demands, and on this stack it fails from the third call onwards, which
 246 disguises the cause. Appendix A diagnoses it (the TensorExpr fuser and e3nn’s scripted spherical
 247 harmonics [Geiger and Smidt, 2022]); Appendix D shows that compilation does not compose with it;
 248 without the fix we would have compared against a Python loop and overstated every speedup below.

249 Two ratios answer different questions and we never conflate them. With $T(L)$ the latency of an
 250 L -lane reverse pass, the *incremental UQ* speedup $[T(M) - T(1)]/[T(1+K) - T(1)]$ is the cost of
 251 uncertainty once the mean force is paid for, while the *total force-plus-UQ* speedup $T(M)/T(1+K)$
 252 includes it. We lead with the total, because the incremental ratio is unstable at small batch precisely
 253 when $T(1+K) \approx T(1)$ — which is the effect this section is about.

254 That effect is large. Batched reverse mode is nearly *flat* in the lane count at $B=1$ (25.5 ms for one
 255 lane, 30.5 ms for eight), so extra cotangents are almost free and the total speedup at $K=3$ collapses
 256 to $1.11\times$ — against $1.83\times$ if one compares only to a serial loop. At $B=16$ the batched path scales
 257 with L again and the speedup recovers to $1.85\times$; by $B=64$ it is slower than the serial loop at *every*
 258 lane count measured ($1.22\times$ at $L=8$), so serial becomes the strongest baseline again and the speedup
 259 is $1.86\times$. Incremental speedups are $1.63\times$ and $3.18\times$ at $B=1$ and $B=16$. Appendix C tabulates the
 260 latencies behind these ratios and Appendix E profiles where the GPU time goes.

261 5.3 Can approximate screening avoid exact computation?

262 This is where the method succeeds — and, unlike Section 5.1, it succeeds under the primary max-
 263 component rule. With a control variate at $r_0=2$, $K=4$ and $\alpha=0.05$, evaluated on the held-out test
 264 split, with Q_{r_0} and τ fitted on the design split and c_α on the calibration split:

system	high-UQ recall	exact skipped	screening speedup
3BPA @ 1200 K	0.972 [0.949, 0.991]	0.844 [0.828, 0.860]	$1.41\times$
rMD17 ethanol	0.982 [0.940, 1.000]	0.836 [0.812, 0.859]	$1.40\times$
rMD17 aspirin	0.973 [0.946, 0.994]	0.851 [0.828, 0.873]	$1.41\times$
rMD17 azobenzene	0.964 [0.926, 0.992]	0.872 [0.849, 0.894]	$1.43\times$

Between 0.836 and 0.872 of exact force-UQ evaluations are avoided while 0.964–0.982 of high-uncertainty configurations are still routed to the exact computation, on every system. Intervals are 95% paired bootstraps over test structures (10,000 resamples); the positive sets hold only tens of structures, so the recall intervals are correspondingly wide and we report them rather than the point estimates alone. The screening gate is where the extreme-value bias of Section 5.1 stops mattering: c_α is a multiplicative calibration constant, so any bias that is monotone in the estimator’s noise is absorbed rather than propagated.

What a screened evaluation costs. Every gate spends $1+K$ lanes per structure: one mean-force lane, which MD needs regardless, plus K uncertainty lanes. On fallback it does *not* restart: the K directions already evaluated are linearly independent inside the r -dimensional centered subspace, and an orthonormal frame spanning them is a recombination of vectors already in hand, so completing the exact basis costs only $r-K$ further passes. With p the fallback fraction, $T_{\text{screen}} = T(1+K) + pT(r-K)$ against $T_{\text{exact}} = T(M)$ — at $K=4$, $M=8$, $T(5) + pT(3)$, while the free energy gate, evaluating no centered direction, pays $T(1) + pT(7)$. Crediting that reuse to the control variate alone would manufacture much of its advantage, so every gate is charged this way. $T(\cdot)$ is the measured median at $B=16$ under whichever implementation is faster at that lane count (Section 5.2); fallback is charged a fresh forward rather than a retained graph, which is conservative; and the quoted figure is the *total* force-plus-UQ speedup.

Is the gate doing the work, or is ForceSketch? A calibrated gate is a generic wrapper — any score correlated with exact force disagreement can be calibrated the same way — so the claim only stands if this gate beats the cheaper ones a reader will reach for. We build three at matched lane budget with identical splits, calibration and threshold. The first is *free*: the M head energies come from the forward pass the model already runs, costing zero *uncertainty* lanes. The second is head subsampling using the *exact* mean force ForceSketch already pays a lane for, $\hat{v}_d = \frac{M}{K(M-1)} \sum_{i \in \mathcal{S}} (F_{di} - \bar{F}_d)^2$, unbiased under uniform sampling and markedly lower-variance than the sample variance among the drawn heads.

gate	lanes	exact skipped	high-UQ recall	speedup
energy disagreement (free)	1	0.294–0.525	0.636–1.000	1.05–1.33×
head subsampling, exact mean	5	0.474–0.747	0.847–0.936	1.16–1.33×
Haar sketch	5	0.525–0.734	0.870–0.926	1.19–1.32×
control variate	5	0.836–0.872	0.964–0.982	1.40–1.43×

Ranges are across the four systems; per-system values with 95% intervals are in the artifact. The control variate *Pareto-dominates* — strictly more skipped *and* strictly higher recall, not a trade — in 11 of the 12 gate-by-system comparisons, including every one against the two sketch baselines. The exception is worth stating plainly: on ethanol the free energy gate attains perfect recall (1.000) against the control variate’s 0.982, while skipping only 0.322 of evaluations against 0.836. Elsewhere that free gate is the informative failure: cheap but coarse, skipping 0.294 on 3BPA and collapsing to 0.636 [0.429, 0.833] recall on azobenzene, so it cannot be trusted as a safety screen — energy disagreement is a genuinely different signal from force disagreement, not a cheap proxy for it. Head subsampling with the exact mean is statistically almost indistinguishable from the Haar sketch at equal budget, as their nearly equal estimator variances predict; what separates the control variate from both is exploiting the concentration in the head-space spectrum rather than probing it obliviously. All three spend identical lanes and all are net speedups at $B=16$, so that margin is decision quality, not cost.

6 Conclusion and limitations

Exact multi-head force uncertainty cannot be replaced by a handful of randomized head-space probes at the fidelity acquisition requires. Across the full budget $K = 1 \dots M-1$, the best sketch reaches only 0.484–0.570 top-5% recall at $K=4$ and 0.724–0.801 even at $K=6$ of 7 under the max-component rule: 0.90 is not attained at *any* budget short of the exact basis. Part of the reason is specific to extreme-value acquisition — the maximum of $3N$ noisy component estimates is biased upward by 1.07–1.08 at $K=4$, which no marginal finite- K correction removes.

It can, however, be *screened*. A split-conformal gate over a control-variate sketch avoids 0.836–0.872 of exact force-UQ computations while retaining 0.964–0.982 of high-uncertainty structures on every system, beating a free energy gate and head subsampling at matched budget in 11 of 12 comparisons. That result is a property of the estimator, not the implementation.

Limitations. Whether it saves *time* is separate, and the honest answer is: only in batch. Once batched reverse mode is available, extra cotangent lanes are nearly free at $B=1$, the total speedup falls to $1.11\times$ and the gate becomes a net slowdown ($0.93\times$); single-structure MD is precisely where this does not pay for itself. The batched advantage is bounded above too ($3.63\times$ at $B=1$, $1.58\times$ at $B=16$, slower than serial by $B=64$), a crossover measured on one 8 GB card without a claimed mechanism. We also preserve a *proxy*, not a demonstrated outcome: we run no active-learning loop and measure no reference-force error, so we claim no gain in model quality or DFT-query efficiency. Further: `torch.func` is unavailable here; committee disagreement is only a proxy for error; calibration is marginal and under exchangeability, not out of distribution; $M=8$ bounds the maximum saving; and rMD17 spans only 9–24 atoms, so molecule size is a weak axis. The method therefore suits batched active-learning sweeps, not single-structure MD steps — and that boundary, with the fact that orthogonalizing the probes costs nothing and buys more than any other estimator choice, is what we think is worth taking away.

Data and code availability

All datasets and model checkpoints used here are public and are pinned by SHA-256 in the artifact’s freeze manifest: the 3BPA [Kovács et al., 2021] and rMD17 [Christensen and von Lilienfeld, 2020] datasets, and the pretrained multi-head committees released with Beck et al. [2025] (Zenodo record 17829635). The estimator library, every experiment script, the analysis code, the ten frozen random seeds and every raw result record behind the numbers in this paper are available at <https://anonymous.4open.science/r/forcesketch-artifact-65F3/README.md>. Everything downstream of the cached head-force matrix is offline linear algebra and reproduces without a GPU; only the timing experiments require one.

References

- Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. In *International Conference on Learning Representations*, 2024. <https://openreview.net/forum?id=33XGfHLtZg>.
- Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, et al. PyTorch 2: Faster machine learning through dynamic Python bytecode transformation and graph compilation. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*, pages 929–947. ACM, 2024. doi:10.1145/3620665.3640366.
- Haim Avron and Sivan Toledo. Randomized algorithms for estimating the trace of an implicit symmetric positive semi-definite matrix. *Journal of the ACM*, 58(2):8:1–8:34, 2011. doi:10.1145/1944345.1944349.
- Killian Bakong, Laurent Massoulié, Edouard Oyallon, and Kevin Scaman. Unbiased approximate vector-jacobian products for efficient backpropagation, 2026. <https://arxiv.org/abs/2602.14701>.
- Ilyes Batatia, Dávid Péter Kovács, Gregor N. C. Simm, Christoph Ortner, and Gábor Csányi. MACE: Higher order equivariant message passing neural networks for fast and accurate force fields. In *Advances in Neural Information Processing Systems 35*, pages 11423–11436, 2022. doi:10.52202/068431-0830.
- Hubert Beck, Pavol Simko, Lars L. Schaaf, Ondrej Marsalek, and Christoph Schran. Multi-head committees enable direct uncertainty prediction for atomistic foundation models. *The Journal of Chemical Physics*, 163(23):234103, 2025. doi:10.1063/5.0302097. arXiv:2508.09907.
- Jörg Behler. Constructing high-dimensional neural network potentials: A tutorial review. *International Journal of Quantum Chemistry*, 115(16):1032–1050, 2015. doi:10.1002/qua.24890.

- 363 Jesús Carrete, Hadrián Montes-Campos, Ralf Wanzenböck, Esther Heid, and Georg K. H. Madsen.
 364 Deep ensembles vs committees for uncertainty estimation in neural-network force fields: Compari-
 365 son and application to active learning. *The Journal of Chemical Physics*, 158(20):204801, 2023.
 366 doi:10.1063/5.0146905.
- 367 Anders S. Christensen and O. Anatole von Lilienfeld. On the role of gradients for machine learning
 368 of molecular energies and forces. *Machine Learning: Science and Technology*, 1(4):045018, 2020.
 369 doi:10.1088/2632-2153/abba6f. rMD17 dataset.
- 370 Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC,
 371 1993. doi:10.1201/9780429246593.
- 372 Mario Geiger and Tess Smidt. e3nn: Euclidean neural networks, 2022. [https://arxiv.org/abs/](https://arxiv.org/abs/2207.09453)
 373 2207.09453.
- 374 Nathan Halko, Per-Gunnar Martinsson, and Joel A. Tropp. Finding structure with randomness:
 375 Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM Review*, 53
 376 (2):217–288, 2011. doi:10.1137/090771806.
- 377 M. F. Hutchinson. A stochastic estimator of the trace of the influence matrix for Laplacian smoothing
 378 splines. *Communications in Statistics — Simulation and Computation*, 19(2):433–450, 1990.
 379 doi:10.1080/03610919008812866.
- 380 Matthias Kellner and Michele Ceriotti. Uncertainty quantification by direct propagation of shallow
 381 ensembles. *Machine Learning: Science and Technology*, 5(3):035006, 2024. doi:10.1088/2632-
 382 2153/ad594a.
- 383 Dávid Péter Kovács, Cas van der Oord, Jiri Kucera, Alice E. A. Allen, Daniel J. Cole, Christoph Or-
 384 tner, and Gábor Csányi. Linear atomic cluster expansion force fields for organic molecules:
 385 Beyond RMSE. *Journal of Chemical Theory and Computation*, 17(12):7696–7711, 2021.
 386 doi:10.1021/acs.jctc.1c00647. 3BPA dataset.
- 387 Raphael A. Meyer, Cameron Musco, Christopher Musco, and David P. Woodruff. Hutch++: Optimal
 388 stochastic trace estimation. In *Symposium on Simplicity in Algorithms (SOSA)*, pages 142–155.
 389 Society for Industrial and Applied Mathematics, 2021. doi:10.1137/1.9781611976496.16.
- 390 Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, et al.
 391 PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural*
 392 *Information Processing Systems 32*, pages 8024–8035, 2019. [https://proceedings.neurips.](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html)
 393 [cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html).
- 394 Evgeny V. Podryabinkin and Alexander V. Shapeev. Active learning of linearly
 395 parametrized interatomic potentials. *Computational Materials Science*, 140:171–180, 2017.
 396 doi:10.1016/j.commatsci.2017.08.031.
- 397 Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random*
 398 *World*. Springer, 2005. doi:10.1007/b106715.
- 399 Felix X. Yu, Ananda Theertha Suresh, Krzysztof Choromanski, Daniel Holtmann-Rice, and Sanjiv
 400 Kumar. Orthogonal random features. In *Advances in Neural Information Processing Systems 29*,
 401 pages 1975–1983, 2016. [https://proceedings.neurips.cc/paper_files/paper/2016/](https://proceedings.neurips.cc/paper_files/paper/2016/hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html)
 402 [hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2016/hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html).

403 A Diagnosing the batched reverse-mode failure

404 **Batched reverse mode, and why it initially appeared impossible.**
 405 `torch.autograd.grad(is_grads_batched=True)` vectorizes the backward over cotan-
 406 gents while reusing one forward graph, and is the baseline spec-level fairness demands. On this
 407 stack it fails with Cannot access data pointer of Tensor that doesn’t have storage
 408 — but only from the third call onwards, which disguises the cause. The mechanism is: batched reverse
 409 mode runs under `torch._vmap_internals`, so cotangents are BatchedTensors with no storage;




Figure 3: Accuracy and cost against K , ranked by the global statistic S . Latency curves coincide for every method at equal lane count under the serial implementation; batched reverse mode flattens them (Section 5.2). Ranking saturates well before tail selection: Haar clears Spearman $\rho = 0.85$ at $K=4$ while its top-5% recall is still 0.619. Error bars are the standard deviation across the ten frozen seeds.

TorchScript’s profiling executor emits an optimized plan only after two warm-ups; and the TensorExpr fuser then fuses the *reverse* graph into a `prim::TensorExprGroup`, whose kernel requires a raw `data_ptr()` for every input. The TorchScript on the path is e3nn’s `_spherical_harmonics`, a module-level `@torch.jit.script free function` — not a module, so walking the module tree finds nothing to un-script, and a fully rebuilt zero-ScriptModule model still fails. Disabling the TensorExpr fuser before the first forward fixes it: 50/50 consecutive batched calls agreeing with the serial path to 3×10^{-15} , at a $\sim 4\%$ cost to the serial path from losing fusion. We report this because it is a trap any e3nn-based UQ work will hit, and because without it we would have compared against a Python loop, which spec-level fairness forbids.

B Accuracy and cost on the global statistic

Figure 2 (left) plots recall against K under the primary max-component rule. For completeness the same trade-off on the easier global statistic S , together with the latency curve, is below; it is the secondary view, and is uniformly more optimistic.

C Measured serial and batched latency

lanes L	$B = 1$ (27 atoms)		$B = 16$ (432 atoms)	
	serial	batched	serial	batched
1	22.9	25.5	24.3	26.0
4	60.3	27.5	62.5	39.7
8	110.5	30.5	115.5	73.2

Median latency in ms, fp32, one RTX 5070 Laptop GPU, CUDA events with one event pair per iteration; **bold** marks the faster implementation at that (B, L) . These are the measurements behind every speedup quoted in Section 5.2.

D Compilation, and why it does not compose

For completeness we also compiled the model. `torch.compile` [Ansel et al., 2024] succeeds only partially here — Dynamo builds a sequence-length guard on e3nn’s `Irrep` and calls `len()` on it, which e3nn raises `NotImplementedError` for by design, so untraceable regions must fall back to eager. With that, compilation is correct to 2.2×10^{-7} and gives $1.04\text{--}1.10\times$ at $B=16$. The absolute saving grows linearly in L , so AOTAutograd is producing a cheaper backward rather than merely a faster forward.

The obvious next question is whether the two accelerations *compose*: compiled-plus-batched would then be the baseline fairness demands. They do not. With the TensorExpr fuser already disabled, compiling the model and then taking a batched VJP reproduces the identical `Cannot access data pointer of Tensor that doesn’t have storage failure`, at both $B=1$ and $B=16$ —

439 AOTAutograd’s compiled backward is no more able to accept a storage-free BatchedTensor than
 440 the fused TorchScript one was. We therefore report the eager batched path as primary because it is
 441 the fastest *correct* baseline available on this stack, not merely the most convenient one, and note that
 442 compilation alone would shift these numbers by at most a few percent in the same direction for every
 443 method, leaving the crossover of Section 5.2 unchanged.

444 E Where the GPU time goes

445 A kernel-level profile of the serial path explains why removing lanes helps at all when it does:
 446 the composition of GPU time is invariant in L (elementwise 42.2% at $L=7$ against 41.5–41.9% at
 447 $L=2, 3, 4$; tensor products 21.3% against 20.0–20.9%; the same 134 distinct kernels throughout),
 448 and only the repetition count changes: kernel launches fall by $2.12\times$ from the exact path to $K=3$,
 449 against a lane ratio of $2.33\times$. Sketching linearly reduces repeated execution of the shared trunk, and
 450 batched reverse mode attacks the same redundancy by a different route — which is exactly why the
 451 two are substitutive at small batch.

452 F Sketching versus using fewer heads

453 At equal total lanes, does a random sketch beat simply computing a few of the head forces exactly?
 454 The comparison hinges on which head-subsampling estimator is used, and the fair one is not the
 455 obvious one. Bare head subsampling takes the sample variance among the K drawn heads. But at
 456 matched budget it also pays a mean-force lane, and can therefore use the *exact* mean:

$$\hat{v}_d^{\text{sub}} = \frac{M}{K(M-1)} \sum_{i \in \mathcal{S}} (F_{di} - \bar{F}_d)^2,$$

457 unbiased under uniform sampling without replacement and markedly lower-variance than the sample
 458 variance, because it does not spend degrees of freedom re-estimating a mean already in hand.

459 Against that fair baseline the ordering *reverses*. Haar at $K=3$ scores -0.040 $[-0.074, -0.008]$ against
 460 exact-mean head subsampling at $K=3$ on 3BPA and -0.070 $[-0.126, -0.016]$ on aspirin — both
 461 significantly *negative*, i.e. head subsampling wins — and the difference is not significant on ethanol
 462 or azobenzene. An earlier version of this comparison scored the baseline with the sample variance
 463 while still charging it the mean lane, and reported Haar ahead; that was an artifact of hobbling the
 464 baseline, and we report the corrected result instead. Head subsampling at $K=4$ with no mean lane
 465 also beats Haar at equal lanes (-0.059 $[-0.093, -0.023]$ on 3BPA), but cannot produce the mean force
 466 MD requires.

467 The paper’s estimator-design claims do not rest on this comparison. Orthogonalization over Gaussian
 468 probes (0.225 $[0.186, 0.263]$, significant on all four systems) and the control variate over Haar (0.111
 469 $[0.080, 0.145]$) are unaffected, and it is the control variate — not the oblivious sketch — that wins
 470 the screening comparison of Section 5.3, where exact-mean head subsampling is included as a gate
 471 baseline in its fair form.