
Force sketching in MD: Can we replace multi-head committees with screening methods for faster force uncertainty compute?

Anonymous Author(s)

Affiliation

Address

email

Abstract

Machine-learned interatomic potentials simulate molecules far more cheaply than quantum chemistry, but a simulator-based workflow only benefits if the emulator can say where it is misspecified — that is what decides when to fall back on the expensive mechanistic model. Committee disagreement is the usual signal, and a *multi-head* committee makes it nearly free for energies: M heads share one trunk, so one forward pass gives all M energies. Forces are not free. Forces are gradients, and each head’s forces need their own backward pass, so force disagreement — the quantity acquisition actually uses — costs about M times more. We ask whether a handful of cheap random probes can estimate it instead, on a pretrained eight-head committee across four molecular systems. The answer is no for *replacement* and yes for *screening*. Even with six of the seven probes the exact computation would use, the best estimator recovers only 72–80% of the genuinely most uncertain structures. But a calibrated filter built from the same score safely skips 84–87% of the expensive computations while still catching 96–98%, beating cheaper alternatives in 11 of 12 comparisons. The wall-clock benefit is regime-dependent: $1.85\times$ in batch, only $1.11\times$ one structure at a time. The decision-quality result is robust; the speed result is not.

1 Introduction

Uncertainty quantification is what makes learned interatomic potentials usable beyond their training distribution: it drives active learning, flags extrapolation during molecular dynamics, and decides when a structure must go back to the quantum-chemistry calculation the model was trained to imitate, orders of magnitude slower and the very thing we are avoiding. Committee disagreement remains the most widely used estimator of that error: where training data pinned the model down its members agree, and where it did not they diverge, so their spread stands in for an error we cannot see [Behler, 2015, Podryabinkin and Shapeev, 2017, Carrete et al., 2023]. Multi-head committees [Beck et al., 2025] — shallow ensembles more broadly [Kellner and Ceriotti, 2024] — make that cheap by attaching M readout heads to one shared message-passing trunk, so all M energies come from a single forward pass instead of M separately trained models.

That amortization does not extend to forces. As the multi-head committee work states directly, “a single backward pass can not yield the forces for all the heads ... the standard deviation requires multiple reverse passes through the computational graph.” Forces are what MD and most acquisition rules actually consume, so the expensive half of committee uncertainty is exactly the half the shared trunk fails to amortize.

This is a question about an imperfect simulator. A learned potential is an emulator of an expensive mechanistic model — density-functional theory — and is misspecified unevenly across configuration

space. A simulator-based workflow does not need it right everywhere, only a trustworthy account of *where* it is not, so the expensive model is called exactly there. That multi-fidelity decision rests on committee disagreement; we ask how cheaply that signal can be had, and where it cannot.

Downstream, though, one rarely needs every member force: acquisition needs a *ranking*, a safety gate needs a *threshold decision*, error analysis needs the most uncertain *atom*. That motivates estimating the force-variance statistic directly rather than reconstructing the member forces defining it. We call this ForceSketch: replace the exact $(M-1)$ -direction centered head basis with K randomized or orthogonal directions, each still costing one reverse pass.

2 Problem formulation

Setup and notation. A structure of N atoms has Cartesian coordinates $x \in \mathbb{R}^D$ with $D = 3N$. We index a single coordinate by $d \in \{1, \dots, D\}$, an atom by $a \in \{1, \dots, N\}$, and the three Cartesian components of one atom by $\gamma \in \{1, 2, 3\}$, so the pair (a, γ) names the same thing as d . The committee has M heads: M small readout networks attached to one shared message-passing trunk. Because they share the trunk, a single forward pass returns all M head energies at once, as the vector $E(x) = [E_1(x), \dots, E_M(x)]^\top \in \mathbb{R}^M$. Head m 's force field is the negative energy gradient $f_m = -\nabla_x E_m \in \mathbb{R}^D$, and we collect the M of them as the columns of

$$F(x) = [f_1 \cdots f_M] \in \mathbb{R}^{D \times M}, \quad \text{so } F_{dm} \text{ is head } m\text{'s force along coordinate } d.$$

Write $\bar{F}_d = \frac{1}{M} \sum_{m=1}^M F_{dm}$ for the committee-mean force on coordinate d — the force the simulation actually uses.

Two disagreement scores, at two scales. Uncertainty is how much the heads disagree. On a *single coordinate* that is their sample variance; summing over coordinates gives one number for the *whole structure*:

$$\underbrace{v_d}_{\text{per coordinate}} = \frac{1}{M-1} \sum_{m=1}^M (F_{dm} - \bar{F}_d)^2, \quad \underbrace{S}_{\text{per structure}} = \sum_{d=1}^D v_d = \frac{\|FP\|_F^2}{M-1}. \quad (1)$$

Here $\mathbf{1} \in \mathbb{R}^M$ is the all-ones vector, $P = I_M - \frac{1}{M} \mathbf{1}\mathbf{1}^\top$ is the projector that subtracts the head mean (*centering*), and $\|\cdot\|_F$ is the Frobenius norm. Centering is what turns raw predictions into disagreement: FP holds each head's deviation from the consensus, and v_d is just the squared length of its d -th row, $v_d = \|(FP)_{d,:}\|_2^2 / (M-1)$. Pooling the three components of one atom gives the two atom-level scores used in the literature, $u_a^{\text{RMS}} = (\frac{1}{3} \sum_\gamma v_{a\gamma})^{1/2}$ and $u_a^{\text{MHC}} = \frac{1}{3} \sum_\gamma \sqrt{v_{a\gamma}}$, where $v_{a\gamma}$ is v_d for the coordinate $d = (a, \gamma)$. Section 4 says which of these the acquisition rule ranks by.

Why energies are cheap and forces are not. The M energies cost one forward pass because the heads share a trunk. Forces are gradients, and reverse mode returns the gradient of *one scalar* per backward pass — we call one such pass a *lane*. Since S depends on F only through FP , whose columns sum to zero, it lives in an $(M-1)$ -dimensional centered subspace with orthonormal basis Q ; differentiating the scalar $q_j^\top E$ costs one lane and returns $g_j = Fq_j$, with $v_d = \frac{1}{M-1} \sum_j g_{j,d}^2$. Exact force uncertainty therefore costs M lanes: $M-1$ centered directions plus the mean force MD needs anyway (Appendix D). *Our question is whether downstream decisions really need all $M-1$ of them.*

3 Method

Gaussian and Haar sketches. Our estimators are the molecular-UQ instance of randomized trace and norm estimation [Hutchinson, 1990, Avron and Toledo, 2011, Meyer et al., 2021]; the orthogonal variant is the head-space analogue of orthogonal random features [Yu et al., 2016], and the control variate plays the role that a low-rank sketch plays in randomized linear algebra [Halko et al., 2011]. Unbiased approximate vector–Jacobian products have also been used to cheapen *training* [Bakong et al., 2026], by sparsifying the backward pass inside the network; we instead compress the *output-head query space* at inference, leaving each reverse pass exact and reducing how many are needed to estimate an existing committee statistic. The idea is to stop asking for all $M-1$ centered directions

79 and instead *probe* the centered subspace with $K \ll M - 1$ random ones, then rescale. Each probe
80 still costs exactly one lane, so the question is how much of the disagreement K probes recover.
81 Concretely, draw $z_k \sim \mathcal{N}(0, I_M)$ and center it, $w_k = Pz_k$; one lane returns $g_k = Fw_k \in \mathbb{R}^D$, and

$$\hat{v}_d = \frac{1}{K(M-1)} \sum_{k=1}^K g_{k,d}^2 \quad (2)$$

82 is an unbiased estimator of v_d — on average it gives the right answer, though any single draw does
83 not. The intuition is the standard one behind randomized trace estimation: a random probe sees each
84 direction of the subspace equally in expectation, so the average of K probes estimates the total energy
85 in all $M - 1$ of them.

86 Random probes can overlap, wasting budget. Sampling them *orthogonally* instead removes that waste.
87 Let $O \in \mathbb{R}^{(M-1) \times K}$ have orthonormal columns drawn Haar-uniformly (that is, uniformly over all
88 K -dimensional orientations, obtained in practice by orthonormalizing a Gaussian matrix), and set
89 $W = QO \in \mathbb{R}^{M \times K}$: a random orthonormal K -frame lying inside the centered subspace. Collecting
90 the K lanes it costs into $G = FW \in \mathbb{R}^{D \times K}$ — so G_{dk} is the response of coordinate d to probe k —
91 gives

$$\hat{v}_d^{\text{Haar}} = \frac{1}{K} \sum_{k=1}^K G_{dk}^2, \quad (3)$$

92 which is unbiased and becomes *exact* at $K = M - 1$, where the frame is a full basis of the centered
93 subspace. Writing $r = M - 1$ for its dimension, the two estimators have relative standard deviations
94 $\sqrt{2/K}$ and $\sqrt{2(r-K)/(K(r+2))}$, so orthogonalization multiplies the estimator variance by
95 exactly $(r-K)/(r+2)$ — a factor of $4/9$ at $K=3, M=8$. It costs nothing extra, and it is the single
96 largest methodological lever we find.

97 **Finite- K corrections.** $\sqrt{\hat{v}_d}$ is biased low, because the square root is concave. A multiplicative
98 constant depending only on K (and, for orthogonal probes, on r) removes that bias; Appendix I gives
99 both constants, including a form of the Haar constant that is finite at $K = r$ where the conventional
100 Beta expression is singular.

101 **Low-rank control variate.** Disagreement is not spread evenly over the centered subspace, so
102 probing it obviously wastes budget. Split $P = Q_{r_0}Q_{r_0}^\top + R$, evaluate the r_0 leading directions of
103 the design-split average $\bar{A} = \frac{1}{n} \sum_i P F_i^\top F_i P$ exactly, and sketch only the residual with $K' = K - r_0$
104 Haar probes. The estimator stays unbiased (Appendix E).

105 **Screening rather than replacement.** If the sketch cannot *replace* the exact score, it can still
106 certify when that score is small. Split conformal [Vovk et al., 2005, Angelopoulos et al., 2024] turns
107 any score into a calibrated bound from held-out data alone: on a calibration split form the ratios
108 $\eta_i = S(x_i)/(\hat{S}(x_i) + \epsilon)$ and let c_α be their $\lceil (n+1)(1-\alpha) \rceil$ -th order statistic. Then $U(x) = c_\alpha \hat{S}(x)$
109 sits above $S(x)$ at least $1 - \alpha$ of the time, so $U(x) < \tau$ certifies the exact score is below the
110 high-uncertainty threshold: skip when it is, fall back otherwise. This is *marginal* coverage under
111 exchangeability, not a conditional or out-of-distribution guarantee (Appendix F).

112 4 Experimental setup

113 We use the pretrained eight-head MACE [Batatia et al., 2022] committees released with Beck
114 et al. [2025]: two message-passing layers, 32 channels, $\ell_{\max} = 1$, $r_{\text{cut}} = 6 \text{ \AA}$. We evaluate the
115 three head-distribution variants (*disjoint*, *overlapping*, *same*) on 3BPA [Kovács et al., 2021] at
116 1200 K (2139 structures, extrapolative relative to the 300 K training set) and the joint ten-molecule
117 rMD17 [Christensen and von Lilienfeld, 2020] committee on ethanol (9 atoms), aspirin (21) and
118 azobenzene (24), 1000 held-out structures each. Estimator evaluation is in float64 and this is not
119 fastidiousness: the mean force exceeds the disagreement by a factor $\kappa \approx 36$ here, so a float32 “exact”
120 reference carries more relative error than the 10^{-5} tolerance at which we verify full-rank exactness.
121 Deployment timing is float32, on one RTX 5070 Laptop GPU, with PyTorch [Paszke et al., 2019]
122 CUDA events measured per iteration; the fitted lane cost is $T(L) = 10.38 + 11.78 L$ ms. Ten random

seeds are frozen before testing and none discarded; all intervals are paired bootstraps [Efron and Tibshirani, 1993] over structures with 10,000 resamples.

The acquisition decision. Following Beck et al. [2025], structures are ranked by the maximum disagreement of their force components. We take the max over components, $\hat{u}^{\max\text{-comp}} = \max_d \hat{\sigma}_d$, as primary and the global sum $\hat{S} = \sum_d \hat{v}_d$ as an easier secondary; a hat always denotes the sketched estimate of the corresponding exact quantity. Section 5.1 measures top-5% recall against the exact score; Section 5.3 applies a fixed threshold τ transferred from the design split, so its realised test prevalence is 0.028–0.050. Every metric is computed per seed and then averaged, so the deployed budget is K lanes, not $10K$ (Appendix G).

Each system is split once into three disjoint parts — *design* (20%, fitting Q_{r_0} and τ), *calibration* (20%, fitting c_α alone) and *test* (60%, touched once). Split-conformal coverage requires the score function to be frozen independently of the calibration observations, so a basis fitted on calibration structures would void the guarantee (Appendix J).

5 Results

5.1 Can force disagreement be sketched?

Partially, and not well enough to replace exact computation — under either acquisition rule, and markedly less well under the primary one.

Under $u^{\max\text{-comp}}$, Haar is the best sketch at every budget, but its top-5% recall is only 0.362–0.489 at $K=3$ and 0.484–0.570 at $K=4$ across the 4 systems. Rank correlation is far healthier than tail selection (Spearman 0.648–0.781 at $K=4$), which is the recurring pattern: a sketch that orders structures tolerably still fails to identify *which* ones land in the extreme tail. On the easier global statistic the same estimator reaches 0.529 [0.490, 0.566] and 0.619 [0.580, 0.648] at $K=3, 4$ on 3BPA — roughly ten points higher, and still far from substitution quality; Appendix N plots that easier view against K alongside the latency curve. Intervals are 95% paired bootstraps over structures (10,000 resamples) on the seed-averaged statistic; seed-to-seed standard deviations are 0.02–0.09 and are recorded separately.

The extreme-value rule is harder for a reason we measure: a maximum over $3N$ noisy estimates is biased upward by 1.07–1.08 at $K=4$, which no marginal finite- K correction removes. That obstructs replacement but not screening, since a scale-calibrated gate absorbs a monotone bias into c_α (Appendix G).

Does sketching beat simply using fewer heads? It depends on whether an exact mean force is required, and the two framings give opposite answers, so Appendix R reports both rather than settling the comparison by definition.

5.2 Does reducing directions reduce GPU time?

Only in some regimes. Against a *batched* exact baseline — which this stack makes available only after disabling the TensorExpr fuser (Appendix M) — the total force-plus-UQ speedup at $K=3$ is $1.85\times$ at $B=16$ but collapses to $1.11\times$ at $B=1$, where extra cotangent lanes are nearly free. Appendix K gives the full accounting, the latency table and a kernel profile.

5.3 Can approximate screening avoid exact computation?

This is where the method succeeds — and, unlike Section 5.1, it succeeds under the primary max-component rule. With a control variate at $r_0=2$, $K=4$ and $\alpha=0.05$, evaluated on the held-out test split, with Q_{r_0} and τ fitted on the design split and c_α on the calibration split:

system	high-UQ recall	exact skipped	screening speedup
3BPA @ 1200 K	0.972 [0.949, 0.991]	0.844 [0.828, 0.860]	$1.41\times$
rMD17 ethanol	0.982 [0.940, 1.000]	0.836 [0.812, 0.859]	$1.40\times$
rMD17 aspirin	0.973 [0.946, 0.994]	0.851 [0.828, 0.873]	$1.41\times$
rMD17 azobenzene	0.964 [0.926, 0.992]	0.872 [0.849, 0.894]	$1.43\times$




Figure 1: The negative result, under the *primary* max-component rule: top-5% recall against budget K on all four systems. Recall climbs steadily but never reaches the 0.90 replacement line short of the exact basis at $K=7$. Error bars are the standard deviation over ten frozen seeds. Appendix B plots the corresponding screening Pareto.

Between 0.836 and 0.872 of exact force-UQ evaluations are avoided while 0.964–0.982 of high-uncertainty configurations are still routed to the exact computation, on every system. Intervals are 95% paired bootstraps over test structures (10,000 resamples); the positive sets hold only tens of structures, so the recall intervals are correspondingly wide and we report them rather than the point estimates alone. The screening gate is where the extreme-value bias of Section 5.1 stops mattering: c_α is a multiplicative calibration constant, so any bias that is monotone in the estimator’s noise is absorbed rather than propagated.

Every gate spends $1+K$ lanes per structure and, on fallback, only $r-K$ further passes rather than a fresh M — charged identically for all gates, since crediting that reuse to one method alone would manufacture its advantage (Appendix L).

Is the gate doing the work, or is ForceSketch? A calibrated gate is a generic wrapper, so the claim only stands if this gate beats the cheaper ones a reader will reach for. Against a *free* energy-disagreement gate, head subsampling with the exact mean, and a plain Haar sketch — all at matched lane budget, identical splits and threshold — the control variate *Pareto-dominates* in 11 of 12 gate-by-system comparisons. The free gate is the informative failure: its recall collapses to 0.636 on azobenzene, so energy disagreement is a genuinely different signal from force disagreement, not a cheap proxy for it. Appendix A gives the full comparison.

6 Conclusion and limitations

Exact multi-head force uncertainty cannot be replaced by a handful of randomized head-space probes: the best sketch reaches only 0.484–0.570 top-5% recall at $K=4$ and 0.724–0.801 even at $K=6$ or 7. It can be *screened*: a split-conformal gate over a control-variate sketch avoids 0.836–0.872 of exact computations while retaining 0.964–0.982 of high-uncertainty structures on every system — a property of the estimator, not the implementation. **Limitations.** Whether it saves *time* depends on the regime: at $B=1$ extra cotangent lanes are nearly free, the speedup falls to $1.11\times$ and the gate becomes a net slowdown, so this suits batched active-learning sweeps, not single-structure MD. We preserve a *proxy*, not a demonstrated outcome: no active-learning loop is run and no reference-force error measured, so we claim no gain in model quality or oracle-query efficiency. Committee disagreement is only a proxy for error; calibration is marginal and under exchangeability; $M=8$ bounds the maximum saving; and rMD17 spans only 9–24 atoms.

References

Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. In *International Conference on Learning Representations*, 2024. <https://openreview>.

198 net/forum?id=33XGfHLtZg.

199 Jason Ansel, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky,
200 et al. PyTorch 2: Faster machine learning through dynamic Python bytecode transformation and
201 graph compilation. In *Proceedings of the 29th ACM International Conference on Architectural*
202 *Support for Programming Languages and Operating Systems, Volume 2*, pages 929–947. ACM,
203 2024. doi:10.1145/3620665.3640366.

204 Haim Avron and Sivan Toledo. Randomized algorithms for estimating the trace of an im-
205 plicit symmetric positive semi-definite matrix. *Journal of the ACM*, 58(2):8:1–8:34, 2011.
206 doi:10.1145/1944345.1944349.

207 Killian Bakong, Laurent Massoulié, Edouard Oyallon, and Kevin Scaman. Unbiased approximate
208 vector-jacobian products for efficient backpropagation, 2026. [https://arxiv.org/abs/2602.](https://arxiv.org/abs/2602.14701)
209 14701.

210 Ilyes Batatia, Dávid Péter Kovács, Gregor N. C. Simm, Christoph Ortner, and Gábor Csányi.
211 MACE: Higher order equivariant message passing neural networks for fast and accurate force
212 fields. In *Advances in Neural Information Processing Systems 35*, pages 11423–11436, 2022.
213 doi:10.52202/068431-0830.

214 Hubert Beck, Pavol Simko, Lars L. Schaaf, Ondrej Marsalek, and Christoph Schran. Multi-head
215 committees enable direct uncertainty prediction for atomistic foundation models. *The Journal of*
216 *Chemical Physics*, 163(23):234103, 2025. doi:10.1063/5.0302097. arXiv:2508.09907.

217 Jörg Behler. Constructing high-dimensional neural network potentials: A tutorial review. *International*
218 *Journal of Quantum Chemistry*, 115(16):1032–1050, 2015. doi:10.1002/qua.24890.

219 Jesús Carrete, Hadrián Montes-Campos, Ralf Wanzemböck, Esther Heid, and Georg K. H. Madsen.
220 Deep ensembles vs committees for uncertainty estimation in neural-network force fields: Compari-
221 son and application to active learning. *The Journal of Chemical Physics*, 158(20):204801, 2023.
222 doi:10.1063/5.0146905.

223 Anders S. Christensen and O. Anatole von Lilienfeld. On the role of gradients for machine learning
224 of molecular energies and forces. *Machine Learning: Science and Technology*, 1(4):045018, 2020.
225 doi:10.1088/2632-2153/abba6f. rMD17 dataset.

226 Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman and Hall/CRC,
227 1993. doi:10.1201/9780429246593.

228 Mario Geiger and Tess Smidt. e3nn: Euclidean neural networks, 2022. [https://arxiv.org/abs/](https://arxiv.org/abs/2207.09453)
229 2207.09453.

230 Nathan Halko, Per-Gunnar Martinsson, and Joel A. Tropp. Finding structure with randomness:
231 Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM Review*, 53
232 (2):217–288, 2011. doi:10.1137/090771806.

233 M. F. Hutchinson. A stochastic estimator of the trace of the influence matrix for Laplacian smoothing
234 splines. *Communications in Statistics — Simulation and Computation*, 19(2):433–450, 1990.
235 doi:10.1080/03610919008812866.

236 Matthias Kellner and Michele Ceriotti. Uncertainty quantification by direct propagation of shallow
237 ensembles. *Machine Learning: Science and Technology*, 5(3):035006, 2024. doi:10.1088/2632-
238 2153/ad594a.

239 Dávid Péter Kovács, Cas van der Oord, Jiri Kucera, Alice E. A. Allen, Daniel J. Cole, Christoph Or-
240 tner, and Gábor Csányi. Linear atomic cluster expansion force fields for organic molecules:
241 Beyond RMSE. *Journal of Chemical Theory and Computation*, 17(12):7696–7711, 2021.
242 doi:10.1021/acs.jctc.1c00647. 3BPA dataset.

243 Raphael A. Meyer, Cameron Musco, Christopher Musco, and David P. Woodruff. Hutch++: Optimal
244 stochastic trace estimation. In *Symposium on Simplicity in Algorithms (SOSA)*, pages 142–155.
245 Society for Industrial and Applied Mathematics, 2021. doi:10.1137/1.9781611976496.16.

- 246 Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, et al.
 247 PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural*
 248 *Information Processing Systems 32*, pages 8024–8035, 2019. [https://proceedings.neurips.](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html)
 249 [cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html).
- 250 Evgeny V. Podryabinkin and Alexander V. Shapeev. Active learning of linearly
 251 parametrized interatomic potentials. *Computational Materials Science*, 140:171–180, 2017.
 252 doi:10.1016/j.commatsci.2017.08.031.
- 253 Vladimir Vovk, Alexander Gammernan, and Glenn Shafer. *Algorithmic Learning in a Random*
 254 *World*. Springer, 2005. doi:10.1007/b106715.
- 255 Felix X. Yu, Ananda Theertha Suresh, Krzysztof Choromanski, Daniel Holtmann-Rice, and Sanjiv
 256 Kumar. Orthogonal random features. In *Advances in Neural Information Processing Systems 29*,
 257 pages 1975–1983, 2016. [https://proceedings.neurips.cc/paper_files/paper/2016/](https://proceedings.neurips.cc/paper_files/paper/2016/hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html)
 258 [hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html](https://proceedings.neurips.cc/paper_files/paper/2016/hash/53adaf494dc89ef7196d73636eb2451b-Abstract.html).

259 **NeurIPS Paper Checklist**

260 **1. Claims**

261 Question: Do the main claims made in the abstract and introduction accurately reflect the
 262 paper’s contributions and scope?

263 Answer: [Yes]

264 Justification: The abstract leads with the negative result (at $K \leq 4$ no estimator recovers the
 265 exact top-5% set) and confines the positive claim to screening. It also gives the batch size at
 266 which the wall-clock benefit disappears.

267 Guidelines:

- 268 • The answer [N/A] means that the abstract and introduction do not include the claims
 269 made in the paper.
- 270 • The abstract and/or introduction should clearly state the claims made, including the
 271 contributions made in the paper and important assumptions and limitations. A [No] or
 272 [N/A] answer to this question will not be perceived well by the reviewers.
- 273 • The claims made should match theoretical and experimental results, and reflect how
 274 much the results can be expected to generalize to other settings.
- 275 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
 276 are not attained by the paper.

277 **2. Limitations**

278 Question: Does the paper discuss the limitations of the work performed by the authors?

279 Answer: [Yes]

280 Justification: Section 6 states that the gate is a net slowdown at batch size 1, that the batched-
 281 versus-serial crossover is specific to our 8 GB GPU, and that we measure a proxy rather than
 282 downstream model quality.

283 Guidelines:

- 284 • The answer [N/A] means that the paper has no limitation while the answer [No] means
 285 that the paper has limitations, but those are not discussed in the paper.
- 286 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 287 • The paper should point out any strong assumptions and how robust the results are to
 288 violations of these assumptions (e.g., independence assumptions, noiseless settings,
 289 model well-specification, asymptotic approximations only holding locally). The authors
 290 should reflect on how these assumptions might be violated in practice and what the
 291 implications would be.
- 292 • The authors should reflect on the scope of the claims made, e.g., if the approach was
 293 only tested on a few datasets or with a few runs. In general, empirical results often
 294 depend on implicit assumptions, which should be articulated.

- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Each estimator identity is stated with its assumptions (centered probes, Haar frames, uniform sampling without replacement) and derived in the appendices. Every one is additionally verified by Monte Carlo at 4×10^5 draws.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Model checkpoints and datasets are public and pinned by SHA-256, together with the environment lockfile, the exact library commit, the three data splits and the ten random seeds, which were frozen before the final evaluation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often

one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The estimator library, experiment scripts, analysis code and every raw result record are available as an anonymised artifact; the datasets and checkpoints are public and identified by content hash in its freeze manifest.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 gives the committee architecture, the systems and splits, the probe budgets K , the control-variate rank r_0 , the miscoverage level α and the precision used for correctness versus timing.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All intervals are 95% paired bootstraps over structures with 10^4 resamples, computed on the seed-averaged statistic; seed-to-seed standard deviations are reported separately.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All timing runs on a single RTX 5070 Laptop GPU (8 GB), measured with CUDA events per iteration. Everything downstream of the cached head-force matrix is offline linear algebra and needs no GPU.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses public datasets and pretrained checkpoints, involves no human subjects or personal data, and conforms with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: This is a measurement study of uncertainty estimators for molecular simulation. Its effect is on the cost and reliability of existing scientific workflows, and we identify no direct societal risk or benefit that would not be speculative.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: We release analysis code and numerical result records only. No pretrained model, generative system or scraped dataset with misuse potential is released.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: MACE, e3nn and PyTorch, the 3BPA and rMD17 datasets, and the pretrained multi-head committees are cited to their original publications and used under their own licenses; each artifact is pinned by content hash.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The estimator library and the raw result records are documented in the artifact README, which states what each script produces and how to reproduce the reported numbers.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

559 Answer: [N/A]
 560 Justification: The paper involves no crowdsourcing and no research with human subjects.
 561 Guidelines:
 562 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
 563 with human subjects.
 564 • Including this information in the supplemental material is fine, but if the main contribu-
 565 tion of the paper involves human subjects, then as much detail as possible should be
 566 included in the main paper.
 567 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
 568 or other labor should be paid at least the minimum wage in the country of the data
 569 collector.

570 **15. Institutional review board (IRB) approvals or equivalent for research with human**
 571 **subjects**
 572 Question: Does the paper describe potential risks incurred by study participants, whether
 573 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
 574 approvals (or an equivalent approval/review based on the requirements of your country or
 575 institution) were obtained?
 576 Answer: [N/A]
 577 Justification: The paper involves no research with human subjects, so no IRB approval was
 578 required.
 579 Guidelines:
 580 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
 581 with human subjects.
 582 • Depending on the country in which research is conducted, IRB approval (or equivalent)
 583 may be required for any human subjects research. If you obtained IRB approval, you
 584 should clearly state this in the paper.
 585 • We recognize that the procedures for this may vary significantly between institutions
 586 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
 587 guidelines for their institution.
 588 • For initial submissions, do not include any information that would break anonymity (if
 589 applicable), such as the institution conducting the review.

590 **16. Declaration of LLM usage**
 591 Question: Does the paper describe the usage of LLMs if it is an important, original, or
 592 non-standard component of the core methods in this research? Note that if the LLM is used
 593 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
 594 scientific rigor, or originality of the research, declaration is not required.
 595 Answer: [Yes]
 596 Justification: An AI coding assistant was used to implement the estimator library and analysis
 597 code, to run experiments, and to draft and revise the manuscript. It is not a component of
 598 the scientific method, and the author directed the design and verified every reported result.
 599 Guidelines:
 600 • The answer [N/A] means that the core method development in this research does not
 601 involve LLMs as any important, original, or non-standard components.
 602 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
 603 be described.

604 **A Screening-gate baselines at matched budget**

605 **Is the gate doing the work, or is ForceSketch?** A calibrated gate is a generic wrapper — any score
 606 correlated with exact force disagreement can be calibrated the same way — so the claim only stands
 607 if this gate beats the cheaper ones a reader will reach for. We build three at matched lane budget
 608 with identical splits, calibration and threshold. The first is *free*: the M head energies come from the

forward pass the model already runs, costing zero *uncertainty* lanes. The second is head subsampling using the *exact* mean force ForceSketch already pays a lane for, $\hat{v}_d = \frac{M}{K(M-1)} \sum_{i \in S} (F_{di} - \bar{F}_d)^2$, unbiased under uniform sampling and markedly lower-variance than the sample variance among the drawn heads.

gate	lanes	exact skipped	high-UQ recall	speedup
energy disagreement (free)	1	0.294–0.525	0.636–1.000	1.05–1.33×
head subsampling, exact mean	5	0.474–0.747	0.847–0.936	1.16–1.33×
Haar sketch	5	0.525–0.734	0.870–0.926	1.19–1.32×
control variate	5	0.836–0.872	0.964–0.982	1.40–1.43×

Ranges are across the four systems. The control variate *Pareto-dominates* — strictly more skipped and strictly higher recall — in 11 of the 12 gate-by-system comparisons, including every one against the two sketch baselines. The exception: on ethanol the free energy gate attains perfect recall (1.000) while skipping only 0.322. Elsewhere that free gate is the informative failure, collapsing to 0.636 recall on azobenzene: energy disagreement is a genuinely different signal from force disagreement, not a cheap proxy for it.

B Screening Pareto at matched budget




Figure 2: All four gates at matched lane budget on all four systems, under the max-component rule. Up-and-right is better on both axes, so the control variate’s dominance is read directly. Bars are 95% paired bootstrap intervals.

C Data and code availability

All datasets and model checkpoints used here are public and are pinned by SHA-256 in the artifact’s freeze manifest: the 3BPA [Kovács et al., 2021] and rMD17 [Christensen and von Lilienfeld, 2020] datasets, and the pretrained multi-head committees released with Beck et al. [2025] (Zenodo record 17829635). The estimator library, every experiment script, the analysis code, the ten frozen random seeds and every raw result record behind the numbers in this paper are available at <https://anonymous.4open.science/r/forcesketch-artifact-65F3/README.md>. Everything downstream of the cached head-force matrix is offline linear algebra and reproduces without a GPU; only the timing experiments require one.

D Lanes and the centered head subspace

Why energies are cheap and forces are not. The M energies cost one forward pass because the heads share a trunk. Forces are gradients, and reverse-mode differentiation returns the gradient of *one scalar* per backward pass. Feeding it the mean energy yields the mean force, not the spread; feeding

634 it each E_m in turn yields all M force fields but costs M backward passes. We call one backward pass
635 a *lane*.

636 The saving grace is that S depends on F only through the centered matrix FP , whose columns
637 sum to zero and therefore span at most $M - 1$ dimensions. Let $Q \in \mathbb{R}^{M \times (M-1)}$ have orthonormal
638 columns $q_1, \dots, q_{M-1}$ spanning exactly that centered subspace:

$$Q^\top Q = I_{M-1}, \quad Q^\top \mathbf{1} = 0, \quad QQ^\top = P.$$

639 Each column is one lane’s worth of work. Differentiating the scalar $q_j^\top E$ — a fixed weighted
640 combination of the head energies — costs a single backward pass and returns

$$g_j = -\nabla_x(q_j^\top E) = Fq_j \in \mathbb{R}^D, \quad \text{and} \quad v_d = \frac{1}{M-1} \sum_{j=1}^{M-1} g_{j,d}^2. \quad (4)$$

641 So the exact variance needs $M - 1$ lanes, plus one more along the uniform direction $\frac{1}{M} \mathbf{1}$ for the
642 mean force the simulation needs anyway: M lanes in total, matching the naive per-head loop. *Our*
643 *question is whether downstream uncertainty decisions really need all $M - 1$ centered lanes.*

644 E The low-rank control variate

645 **Low-rank control variate.** Split $P = Q_{r_0} Q_{r_0}^\top + R$, where $Q_{r_0} \in \mathbb{R}^{M \times r_0}$ has orthonormal columns
646 inside the centered subspace and R is the projector onto what is left of it, the *residual* subspace of
647 dimension $r - r_0$. The columns of Q_{r_0} are the r_0 leading eigenvectors of the *design-split* average
648 $\bar{A} = \frac{1}{n} \sum_i P F_i^\top F_i P$ — an $M \times M$ matrix in *head* space, whose top eigenvector is the weighting
649 of heads along which this committee disagrees most consistently. It is fitted separately for each
650 system on structures disjoint from both the calibration and the test splits (Section 4). Fitting it on
651 the calibration split instead would invalidate the conformal guarantee below, since split conformal
652 requires the score function — learned basis included — to be frozen before the calibration ratios are
653 formed. The idea is to spend a few lanes where the disagreement actually concentrates, and probe
654 only what is left over. Evaluate those r_0 leading directions exactly, costing $g_j = Fq_j$ for the columns
655 q_j of Q_{r_0} , and sketch only the residual with $K' = K - r_0$ Haar probes drawn inside R ’s range,
656 collected as $G \in \mathbb{R}^{D \times K'}$ exactly as before but with the frame now living in the residual subspace:

$$\hat{v}_d = \frac{1}{r} \left[\sum_{j \leq r_0} g_{j,d}^2 + \frac{r-r_0}{K'} \sum_{k=1}^{K'} G_{dk}^2 \right], \quad (5)$$

657 which remains unbiased. By *spectrum* we mean the eigenvalues of $\bar{A}$ — how the total disagreement
658 is divided among head-space directions. The motivation is then the *opposite* of the intuitive one: a
659 concentrated spectrum does not explain why small K works, it is what *breaks* a random sketch, since
660 a probe either catches the dominant direction or misses it.

661 F Split-conformal screening

662 **Screening rather than replacement.** If the sketch cannot *replace* the exact score, it can still say
663 when the exact score is certainly small. Conformal prediction [Vovk et al., 2005, Angelopoulos
664 et al., 2024] turns any score, however crude, into a calibrated bound using held-out data alone. On a
665 calibration split of n structures disjoint from design and test, form the ratios $\eta_i = S(x_i)/(\hat{S}(x_i) + \epsilon)$,
666 the factor by which the cheap score underestimates the exact one — and let c_α be their $\lceil (n+1)(1-\alpha) \rceil$ -
667 th order statistic, large enough to cover all but an α fraction of them. Then $U(x) = c_\alpha \hat{S}(x)$ sits
668 above $S(x)$ at least $1 - \alpha$ of the time, so $U(x) < \tau$ *certifies* that the exact score is below the
669 high-uncertainty threshold τ too: skip when $U(x) < \tau$, fall back otherwise. Formally this is *marginal*
670 coverage, $\mathbb{P}[S(x) \leq U(x)] \geq 1 - \alpha$, under exchangeability — not a conditional guarantee and not
671 an out-of-distribution one. 3BPA at 1200 K calibrated against 300 K behaviour is exactly such a shift,
672 so we report the observed coverage there rather than assume it.

673 G The acquisition decision and the extreme-value bias

674 **The acquisition decision, stated exactly.** Every number below concerns one decision: from a
675 pool of held-out structures, select the most force-uncertain fraction for exact reference computation.

Following Beck et al. [2025], structures are ranked by the maximum disagreement of their force components; we evaluate that extreme-value rule in both of its defensible forms — maximum over components, maximum over atoms — and report the global sum as an easier diagnostic statistic:

$$\hat{u}^{\max\text{-comp}}(x) = \max_d \hat{\sigma}_d, \quad \hat{u}^{\max\text{-atom}}(x) = \max_a \hat{u}_a^{\text{MHC}}, \quad \hat{S}(x) = \sum_d \hat{v}_d. \quad (6)$$

A hat always denotes the sketched estimate of the corresponding exact quantity: the exact scores are the same three expressions with v_d and $\sigma_d = \sqrt{v_d}$ in place of $\hat{v}_d$ and $\hat{\sigma}_d$, so the exact S here is exactly the S of Eq. (1). The correction c is c_K for Gaussian probes and $c_{K,r}^{\text{Haar}}$ for orthogonal ones (Section 3); u_a^{MHC} is as in Section 2. We take $u^{\max\text{-comp}}$ as primary, being the most literal reading, report $u^{\max\text{-atom}}$ alongside it in the artifact, and treat S as the statistically *easier* secondary: a sum over $3N$ coordinates concentrates, whereas both maxima are extreme-value statistics of a noisy $\hat{v}$.

The two experiments select positives differently and we do not conflate them. Section 5.1 measures **top-5% recall**: the 5% of evaluated structures with the largest *exact* score, ties resolved inclusively, so the set may slightly exceed 5%. Section 5.3 applies a *fixed* threshold instead, as a deployed gate must: τ is the 95th percentile of the exact score on the design split, and a test structure counts as high-uncertainty when its exact score is at least τ — so the realised test prevalence is 0.028–0.050, not exactly 0.05. Every metric is computed *per seed and then averaged* over the ten frozen seeds: the deployed budget is K lanes, not $10K$, and averaging the sketches before scoring would be a different, far more accurate estimator than the one we claim.

Why the extreme-value rule is harder. Making each $\hat{\sigma}_d$ unbiased *marginally* is not the same as making $\max_d \hat{\sigma}_d$ unbiased: a maximum over $3N$ noisy estimates is biased upward, and the bias grows with the estimator’s noise. Measured directly, $u^{\max\text{-comp}}$ is inflated by 1.42–1.53 at $K=1$, falling to 1.07–1.08 at $K=4$ and to 1.000 at the full basis, where the estimator is exact and the bias must vanish — a useful check on the pipeline. This obstructs *replacement* but not screening (Section 5.3): the bias is monotone in the estimator’s noise, so a scale-calibrated gate absorbs it into c_α .

Orthogonalizing the probes is the largest single lever in the method. The estimator comparisons that follow are on the global statistic, where the paired-bootstrap machinery of Section 4 was frozen; the ordering they establish is a property of the estimators, not of the acquisition rule. The paired difference between Haar and Gaussian at $K=3$, evaluated on identical structure resamples, is 0.225 [0.186, 0.263] in top-5% recall, and is significant on all four systems. The low-rank control variate adds a further 0.111 [0.080, 0.145] over Haar at $K=4$, reaching 0.730 [0.692, 0.760] — still short of 0.90.

The head-space spectrum explains why. Writing $A = Q^T F^T F Q$ for the averaged centered force-disagreement Gram matrix, the stable rank of the matrix we actually sketch is $\text{srank}(FQ) = \|FQ\|_F^2 / \|FQ\|_2^2 = \text{tr}(A) / \lambda_{\max} = 2.91\text{--}3.08$ of a maximum 7 across the four systems, with the leading direction carrying 0.325–0.344 of the disagreement (an isotropic spectrum would give 0.143). The spectrum is therefore *moderately diffuse rather than strongly rank-concentrated* — but far from isotropic, at $2.3\text{--}2.4\times$ the isotropic leading share. That is exactly the regime in which a random K -dimensional projection pays a variance penalty, and it is why projecting out the leading directions (the control variate) recovers so much.

714 H What ForceSketch replaces

715 I Finite- K corrections

Finite- K corrections. Reporting a standard deviation means taking a square root, and the square root of an unbiased variance estimate is *not* an unbiased standard-deviation estimate: $\sqrt{\cdot}$ is concave, so $\sqrt{\hat{v}_d}$ is biased low, the more so the fewer probes we spend. The fix is a multiplicative constant depending only on K (and, for orthogonal probes, on r). For Gaussian probes $K\hat{v}_d/v_d \sim \chi_K^2$ and the correction is $c_K = \sqrt{2/K} \Gamma(\frac{K+1}{2}) / \Gamma(\frac{K}{2})$. For Haar probes the analogous constant is conventionally written with Beta functions, which is singular at $K = r$ (it evaluates $B(K/2, 0)$). Cancelling the shared $\Gamma(\frac{r-K}{2})$ gives a form that is finite everywhere and exactly 1 at full rank:

$$c_{K,r}^{\text{Haar}} = \sqrt{\frac{r}{K}} \frac{\Gamma(\frac{K+1}{2}) \Gamma(\frac{r}{2})}{\Gamma(\frac{K}{2}) \Gamma(\frac{r+1}{2})}, \quad c_{K,r}^{\text{Haar}} \rightarrow c_K \text{ as } r \rightarrow \infty. \quad (7)$$




Figure 3: What ForceSketch replaces. All M head energies come from one forward pass, but forces do not: exact disagreement needs the full centered head subspace, M lanes. ForceSketch spends $K+1$, and — since that is not accurate enough to *replace* exact uncertainty at $K \leq 4$ under either acquisition rule (Section 5.1) — its score feeds a calibrated gate that runs the exact path only where the bound cannot rule out the threshold. Whether fewer lanes also cost less time depends on the autodiff regime (Section K.1).

We verify unbiasedness of $\sqrt{\hat{v}}/c$ to within Monte-Carlo error at 4×10^5 draws. These closed forms cover the Gaussian and Haar estimators; the control variate below mixes exact and sketched directions and has no such constant, so we apply the residual frame’s Haar correction, which is an approximation. It does not affect the screening results, because c_α is itself a multiplicative calibration and absorbs any constant scale error exactly.

J Splits

Splits. Each system is partitioned once, under a frozen permutation, into three disjoint splits: a *design* split (20%) fitting the control-variate basis Q_{r_0} and the threshold τ ; a *calibration* split (20%) fitting the conformal constant c_α alone; and a *test* split (60%), touched once — 427/428/1284 structures for 3BPA and 200/200/600 per rMD17 molecule. The separation is not cosmetic: split-conformal coverage requires the score function to be frozen independently of the calibration observations, so a basis fitted on calibration structures would void the guarantee. The code asserts all three pairwise disjointness conditions and that the splits partition the set.

K Does reducing directions reduce GPU time?

K.1 Does reducing directions reduce GPU time?

Only in some regimes, and the honest answer required first making the strongest baseline available. Batched reverse mode — vectorizing the backward over cotangents while reusing one forward graph — is the baseline fairness demands, and on this stack it fails from the third call onwards, which disguises the cause. Appendix M diagnoses it (the TensorExpr fuser and e3nn’s scripted spherical harmonics [Geiger and Smidt, 2022]); Appendix P shows that compilation does not compose with it; without the fix we would have compared against a Python loop and overstated every speedup below.

Two ratios answer different questions and we never conflate them. With $T(L)$ the latency of an L -lane reverse pass, the *incremental UQ* speedup $[T(M) - T(1)]/[T(1+K) - T(1)]$ is the cost of uncertainty once the mean force is paid for, while the *total force-plus-UQ* speedup $T(M)/T(1+K)$ includes it. We lead with the total, because the incremental ratio is unstable at small batch precisely when $T(1+K) \approx T(1)$ — which is the effect this section is about.

That effect is large. Batched reverse mode is nearly *flat* in the lane count at $B=1$ (25.5 ms for one lane, 30.5 ms for eight), so extra cotangents are almost free and the total speedup at $K=3$ collapses to $1.11\times$ — against $1.83\times$ if one compares only to a serial loop. At $B=16$ the batched path scales with L again and the speedup recovers to $1.85\times$; by $B=64$ it is slower than the serial loop at *every* lane count measured ($1.22\times$ at $L=8$), so serial becomes the strongest baseline again and the speedup is $1.86\times$. Incremental speedups are $1.63\times$ and $3.18\times$ at $B=1$ and $B=16$. Appendix O tabulates the latencies behind these ratios and Appendix Q profiles where the GPU time goes.

L What a screened evaluation costs

What a screened evaluation costs. Every gate spends $1+K$ lanes per structure: one mean-force lane, which MD needs regardless, plus K uncertainty lanes. On fallback it does *not* restart: the K directions already evaluated are linearly independent inside the r -dimensional centered subspace, and an orthonormal frame spanning them is a recombination of vectors already in hand, so completing the exact basis costs only $r-K$ further passes. With p the fallback fraction, $T_{\text{screen}} = T(1+K) + pT(r-K)$ against $T_{\text{exact}} = T(M)$ — at $K=4$, $M=8$, $T(5) + pT(3)$, while the free energy gate, evaluating no centered direction, pays $T(1) + pT(7)$. Crediting that reuse to the control variate alone would manufacture much of its advantage, so every gate is charged this way. $T(\cdot)$ is the measured median at $B=16$ under whichever implementation is faster at that lane count (Section K.1); fallback is charged a fresh forward rather than a retained graph, which is conservative; and the quoted figure is the *total* force-plus-UQ speedup.

M Diagnosing the batched reverse-mode failure

Batched reverse mode, and why it initially appeared impossible. `torch.autograd.grad(is_grads_batched=True)` vectorizes the backward over cotangents while reusing one forward graph, and is the baseline spec-level fairness demands. On this stack it fails with Cannot access data pointer of Tensor that doesn't have storage — but only from the third call onwards, which disguises the cause. The mechanism is: batched reverse mode runs under `torch._vmap_internals`, so cotangents are BatchedTensors with no storage; TorchScript's profiling executor emits an optimized plan only after two warm-ups; and the TensorExpr fuser then fuses the *reverse* graph into a `prim::TensorExprGroup`, whose kernel requires a raw `data_ptr()` for every input. The TorchScript on the path is `e3nn's _spherical_harmonics`, a module-level `@torch.jit.script free function` — not a module, so walking the module tree finds nothing to un-script, and a fully rebuilt zero-ScriptModule model still fails. Disabling the TensorExpr fuser before the first forward fixes it: 50/50 consecutive batched calls agreeing with the serial path to 3×10^{-15} , at a $\sim 4\%$ cost to the serial path from losing fusion. We report this because it is a trap any `e3nn`-based UQ work will hit, and because without it we would have compared against a Python loop, which spec-level fairness forbids.

N Accuracy and cost on the global statistic

Figure 1 (left) plots recall against K under the primary max-component rule. For completeness the same trade-off on the easier global statistic S , together with the latency curve, is below; it is the secondary view, and is uniformly more optimistic.

O Measured serial and batched latency

lanes L	$B = 1$ (27 atoms)		$B = 16$ (432 atoms)	
	serial	batched	serial	batched
1	22.9	25.5	24.3	26.0
4	60.3	27.5	62.5	39.7
8	110.5	30.5	115.5	73.2

Median latency in ms, fp32, one RTX 5070 Laptop GPU, CUDA events with one event pair per iteration; **bold** marks the faster implementation at that (B, L) . These are the measurements behind every speedup quoted in Section K.1.




Figure 4: Accuracy and cost against K , ranked by the global statistic S . Latency curves coincide for every method at equal lane count under the serial implementation; batched reverse mode flattens them (Section K.1). Ranking saturates well before tail selection: Haar clears Spearman $\rho = 0.85$ at $K=4$ while its top-5% recall is still 0.619. Error bars are the standard deviation across the ten frozen seeds.

793 P Compilation, and why it does not compose

794 For completeness we also compiled the model. `torch.compile` [Ansel et al., 2024] succeeds only
 795 partially here — Dynamo builds a sequence-length guard on `e3nn`’s `Irrep` and calls `len()` on it,
 796 which `e3nn` raises `NotImplementedError` for by design, so untraceable regions must fall back to
 797 eager. With that, compilation is correct to 2.2×10^{-7} and gives $1.04\text{--}1.10\times$ at $B=16$. The absolute
 798 saving grows linearly in L , so AOTAutograd is producing a cheaper backward rather than merely a
 799 faster forward.

800 The obvious next question is whether the two accelerations *compose*: compiled-plus-batched would
 801 then be the baseline fairness demands. They do not. With the TensorExpr fuser already disabled,
 802 compiling the model and then taking a batched VJP reproduces the identical `Cannot access`
 803 `data pointer of Tensor that doesn’t have storage failure`, at both $B=1$ and $B=16$ —
 804 AOTAutograd’s compiled backward is no more able to accept a storage-free `BatchedTensor` than
 805 the fused TorchScript one was. We therefore report the eager batched path as primary because it is
 806 the fastest *correct* baseline available on this stack, not merely the most convenient one, and note that
 807 compilation alone would shift these numbers by at most a few percent in the same direction for every
 808 method, leaving the crossover of Section K.1 unchanged.

809 Q Where the GPU time goes

810 A kernel-level profile of the serial path explains why removing lanes helps at all when it does:
 811 the composition of GPU time is invariant in L (elementwise 42.2% at $L=7$ against 41.5–41.9% at
 812 $L=2, 3, 4$; tensor products 21.3% against 20.0–20.9%; the same 134 distinct kernels throughout),
 813 and only the repetition count changes: kernel launches fall by $2.12\times$ from the exact path to $K=3$,
 814 against a lane ratio of $2.33\times$. Sketching linearly reduces repeated execution of the shared trunk, and
 815 batched reverse mode attacks the same redundancy by a different route — which is exactly why the
 816 two are substitutive at small batch.

817 R Sketching versus using fewer heads

818 At equal total lanes, does a random sketch beat simply computing a few of the head forces exactly?
 819 The comparison hinges on which head-subsampling estimator is used, and the fair one is not the
 820 obvious one. Bare head subsampling takes the sample variance among the K drawn heads. But at
 821 matched budget it also pays a mean-force lane, and can therefore use the *exact* mean:

$$\hat{v}_d^{\text{sub}} = \frac{M}{K(M-1)} \sum_{i \in S} (F_{di} - \bar{F}_d)^2,$$

822 unbiased under uniform sampling without replacement and markedly lower-variance than the sample
 823 variance, because it does not spend degrees of freedom re-estimating a mean already in hand.

824 Against that fair baseline the ordering *reverses*. Haar at $K=3$ scores -0.040 $[-0.074, -0.008]$ against
825 exact-mean head subsampling at $K=3$ on 3BPA and -0.070 $[-0.126, -0.016]$ on aspirin — both
826 significantly *negative*, i.e. head subsampling wins — and the difference is not significant on ethanol
827 or azobenzene. An earlier version of this comparison scored the baseline with the sample variance
828 while still charging it the mean lane, and reported Haar ahead; that was an artifact of hobbling the
829 baseline, and we report the corrected result instead. Head subsampling at $K=4$ with no mean lane
830 also beats Haar at equal lanes (-0.059 $[-0.093, -0.023]$ on 3BPA), but cannot produce the mean force
831 MD requires.

832 The paper’s estimator-design claims do not rest on this comparison. Orthogonalization over Gaussian
833 probes (0.225 $[0.186, 0.263]$, significant on all four systems) and the control variate over Haar (0.111
834 $[0.080, 0.145]$) are unaffected, and it is the control variate — not the oblivious sketch — that wins
835 the screening comparison of Section 5.3, where exact-mean head subsampling is included as a gate
836 baseline in its fair form.