

WHEN IS COMMITTEE FORCE UNCERTAINTY WORTH COMPUTING IN MD?

DECISION QUALITY, CHEAP GATES, AND THE BASELINES THAT MUST BE BEATEN

Anonymous authors

Paper under double-blind review

ABSTRACT

Molecular dynamics with a machine-learned potential requires repeatedly choosing which structures to label with expensive quantum chemistry. The usual signal is disagreement among an ensemble of models, and cheap approximations of that signal are validated by how closely they match the full ensemble. We instead evaluate the ensemble against the model’s actual error, and evaluate cheap approximations of it separately. Three results follow.

The acquisition rule in common use sums over coordinates, so it grows with molecule size whether or not the model is uncertain. It is harmless at fixed size and misleading when sizes vary, where it selects large structures rather than difficult ones. A deterministic choice of probe directions beats random sampling, skipping 82–94% of the expensive work without changing a decision. On a realistic materials dataset the magnitude of the predicted forces—free, already computed—identifies the worst structures as well as the ensemble does, and the ensemble adds nothing significant on top of it. Ensemble uncertainty is unnecessary where that free signal suffices, and is the baseline future methods should be measured against.

1 INTRODUCTION

Machine-learned interatomic potentials (MLIPs) are trained on labels that cost orders of magnitude more than the model does to evaluate, so which structures get labelled matters more than almost any modelling choice. The dominant selection heuristic is query-by-committee (Seung et al., 1992): train an ensemble, evaluate the spread of its predictions, and label where the spread is largest. For force-field models the natural spread is over *forces*, and obtaining it exactly costs one reverse-mode pass per ensemble member.

A multi-head committee makes that cheaper in principle (Beck et al., 2025). Instead of M independent models, one trunk carries M output heads, so a single reverse pass through the trunk with an appropriate head-space cotangent yields a projection of the per-head force matrix. Whether a small number of such projections suffices to reproduce the acquisition decisions of the exact committee is the question that a sketching analysis is usually asked to answer.

We argue that this is the second question, and that asking it first has hidden a weakness in the whole approach. The literature validates a cheap estimator against the *exact committee*, treating the committee as ground truth. But the committee is itself a proxy—for the model’s error against the reference method—and it is that error a practitioner wants to find. So we decompose:

- **Factor A, oracle quality.** Does exact committee force disagreement rank reference-force error? This is answerable offline from cached per-head forces and requires no sketching at all.
- **Factor B, preservation.** Given whichever oracle one has chosen, how many of its decisions does a cheap gate reproduce, and at what cost?

Factor B is logically independent of Factor A: *whatever you decided to trust, we reproduce it more cheaply* is a claim that survives Factor A being mediocre. Reporting them together, rather than

reporting only the sketch fidelity, is the paper’s main methodological contribution, and it is what surfaces the result in Section 8—that on a realistic materials pool the committee loses to a signal that costs nothing.

Every numeric claim below is derived programmatically from a committed record; none is transcribed. Section 10 describes the reproduction path.

2 RELATED WORK

Committee uncertainty for atomistic models. Query-by-committee (Seung et al., 1992) is the standard selection rule for interatomic potentials, and the standard critique of force-field evaluation is that aggregate error metrics hide the failures that matter (Kovács et al., 2021). Christensen & von Lilienfeld (2020) show that gradient information changes what a model learns, which is part of why force disagreement rather than energy disagreement is the quantity acquisition consumes. Multi-head committees (Beck et al., 2025) amortise the energy spread across one trunk and are the construction we take as given; our contribution is not a better estimator of their statistic but a different question about it. Where that work validates a cheap approximation against the exact committee, we ask first whether the exact committee is worth approximating.

Estimating a trace from few products. The centred head-space statistic is a trace, so the natural tools are randomised. Meyer et al. (2021) establish that splitting a fixed budget between an exact leading subspace and a sketched residual is optimal for trace estimation, and our control variate is that construction. Our negative result is not a contradiction of it: their guarantee is on the trace’s *value*, while acquisition consumes a *ranking* across structures, and Section 6 shows the finite-sample noise that is harmless for the former is decisive for the latter.

Distribution-free calibration. The gate’s constant is the split-conformal one (Vovk et al., 2005; Lei et al., 2018), which buys a finite-sample guarantee under exchangeability alone. Exchangeability is the assumption a molecular dynamics trajectory violates, which is why every split here is blocked with guard bands sized from a measured autocorrelation time rather than drawn at random.

Foundation models and real pools. Universal potentials trained on materials trajectories (Deng et al., 2023; Batatia et al., 2023) are the setting the method is pitched at, and the setting in which we find its motivating premise weakest: a pool spanning composition and size is exactly where an extensive statistic misbehaves and where a free force-magnitude signal is already close to saturated.

3 METHOD

3.1 CENTRED HEAD SPACE

Let $F(x) \in \mathbb{R}^{3N \times M}$ hold the forces predicted by M committee heads for a structure x with N atoms. Acquisition scores depend only on the spread across heads, so the mean-force direction is separated out: with Q_c an $M \times r$ orthonormal (Helmert) basis of the centred subspace, $r = M - 1$, the quantity of interest is $A(x) = Q_c^\top F(x)^\top F(x) Q_c$, an $r \times r$ matrix independent of N . The per-coordinate variance is $v_d = \|Q_c^\top F_d\|^2 / (M - 1)$.

3.2 THREE ACQUISITION STATISTICS

Prior work ranks structures by the largest per-coordinate standard deviation, $\max_d \sqrt{v_d}$ (*max-component*). We also consider the largest per-atom value (*max-atom*) and the sum $\sum_d v_d$ (*global*). The global statistic is a trace of $A(x)$ and so is by far the easiest to estimate from few projections; whether it is also the better physical predictor was pre-registered as a secondary hypothesis before any of it was run.

3.3 ESTIMATORS

A projection (*lane*) costs one reverse pass and returns $g_k = F w_k$ for a head-space direction w_k . We compare Haar and Gaussian sketches, head subsampling, a low-rank control variate that evaluates




Figure 1: Exact committee disagreement ranks reference-force error, and the global statistic ranks it best at fixed system size. AUROC for detecting the top-5% highest-error structures, by signal, on the six molecular systems. The dashed line is the random null; every signal carries a distinct marker so the comparison survives grayscale printing. The two *free* signals cost no backward pass, and the margin over the better of them is 0.067–0.182 AUROC.

r_0 leading directions exactly and sketches the residual—the split that Meyer et al. (2021) show is optimal for trace estimation—and a purely deterministic *leading-only* estimator that spends every lane on exact leading eigendirections of the design-set Gram and discards the residual entirely.

3.4 THE CALIBRATED GATE

The gate skips the exact computation when a cheap score $\hat{S}$ times a calibration constant falls below a threshold. The constant is the split-conformal one: c_α is the $\lceil (n+1)(1-\alpha) \rceil$ -th order statistic of the calibration ratios $S/\hat{S}$ (Vovk et al., 2005; Lei et al., 2018). When $\lceil (n+1)(1-\alpha) \rceil > n$ no order statistic carries the finite-sample guarantee and the implementation returns $+\infty$, which makes the gate skip nothing—the safe failure. Realised coverage is verified against the Beta($k, n+1-k$) law by simulation rather than asserted, at the smallest and largest calibration sets ($n = 84$ and $n = 428$) among the 28 registered splits. The guarantee is weakest at the smallest, $n = 84$, and there realised coverage is 0.952 against the 0.953 the law predicts; 0 of 28 splits lack a feasible order statistic at the α we ship.

What the gate costs when it does *not* skip is worth stating explicitly, because it is the obvious way to flatter such a method. There is no partial credit: a structure the gate declines to skip is evaluated exactly and in full, so the $K+1$ lanes already spent on its cheap score are pure overhead rather than a down payment on the exact basis. The skip fractions below are therefore counts of structures that avoided the exact computation altogether, and the estimators are compared at matched lane budgets.

Design, calibration and test roles are disjoint contiguous blocks with guard bands sized from a measured integrated autocorrelation time; every split is content-hashed and recomputes from its recorded seed. Confidence intervals use a paired moving-block bootstrap (Künsch, 1989).

4 DOES COMMITTEE DISAGREEMENT PREDICT REFERENCE-FORCE ERROR?

Across six molecular systems—3BPA (Kovács et al., 2021) at 1200 K under three committee constructions, and three rMD17 molecules (Christensen & von Lilienfeld, 2020)—exact global disagreement ranks the top-5% highest-error structures at AUROC 0.737–0.848, with a random null at chance. The oracle is *moderate*: useful, and far from perfect. Nothing is fitted in this analysis, so there is no leakage risk.




Figure 2: Randomisation does not earn its lanes. Difference in the fraction of exact evaluations skipped between the deterministic leading-only estimator and the randomised control variate, at three matched lane budgets, with 95% paired block-bootstrap intervals. Positive is leading-only better. It wins significantly on 5 of 6 systems at every budget; the single exception is 3BPA overlapping, in a regime where both estimators already skip more than 93%.

4.1 THE GLOBAL STATISTIC WINS—AT FIXED SYSTEM SIZE

The global statistic beats max-component in 24 of 24 (system $\times$ error-score) cells, 19 of them significantly at 95%, and in no cell is it significantly worse (Figure 1). Because the global statistic is also the one a sketch estimates most easily, the negative result about small sketches (Section 6) applies to the statistic one should not have been using.

This holds *at fixed N* , and the qualifier is load-bearing rather than cosmetic; Section 8.2 shows the ordering reverse when system size varies, and why it could not have shown up here.

4.2 THE MARGIN OVER FREE SIGNALS

Head-energy spread costs one forward pass; the norm of the mean force costs nothing at all. Both are informative (AUROC 0.60–0.70). Exact global disagreement beats the better of them by 0.067–0.182 AUROC, and pairwise the advantage is significant in 9 of 10 comparisons. So the backward passes buy something real but not overwhelming—and Section 8 shows a regime where they buy nothing.

5 WHAT THE CHEAP GATE COSTS

Evaluated on held-out test blocks, the gate skips 0.822–0.943 of exact evaluations while reproducing the exact policy’s decisions. We report this count rather than a speedup because it is independent of hardware, autograd backend and batching.

5.1 RANDOMISATION DOES NOT EARN ITS LANES

At matched lane budget the deterministic leading-only estimator beats the randomised control variate on 5 of 6 systems, at every budget tested, and has zero probe-seed variance by construction (Figure 2). The mechanism is specific to calibrated gates: c_α is fitted to the ratio $S/\hat{S}$, so a systematic bias in $\hat{S}$ is absorbed exactly, and only its *variance* costs skip fraction. Leading-only is biased low by construction and pays nothing for it. Since the head-space spectrum is concentrated—stable rank ≈ 1.2 –4.3 of a possible 7—a random projection is being asked to do the one thing concentration makes hard.

The practical consequence is that the method gets simpler: at fixed budget, spend every lane on exact leading directions.




Figure 3: The finite-sample law predicts the negative result with no free parameters. Predicted against measured top-5% recall for six systems at three sketch budgets. Using the real head-space geometry gives a mean absolute error of 0.015; the isotropic null is both worse and systematically optimistic, which is why coordinate coupling cannot be neglected.

6 THEORY: THE NEGATIVE RESULT IS PREDICTABLE

For a Haar sketch of K of the r centred directions, the ratio of estimated to true per-coordinate variance is distributed exactly as $\hat{v}_d/v_d \sim (r/K) \text{Beta}(K/2, (r-K)/2)$. Injecting precisely this noise into the exact variances and recomputing top-5% recall reproduces the measured recall across six systems and three budgets to a mean absolute error of 0.015, with no free parameters (Figure 3).

Two corrections to the mechanism previously offered follow. First, a common multiplicative bias cannot change a rank statistic, so the damage is estimator *variance*, not bias. Second, the coordinate directions cannot be treated as isotropic: using the real head-space geometry roughly halves the prediction error relative to an isotropic null, so coordinate coupling matters.

A Chernoff bound makes the result structural rather than incidental: uniform componentwise accuracy at $\varepsilon = 0.1$ requires K of order 10^3 – 10^4 , against an available budget of $K \leq r = 7$. The gap is two to three orders of magnitude, so no amount of tuning recovers the max-component rule from a small sketch.

7 GENERALITY

7.1 COMMITTEE SIZE

Enumerating *all* 247 subcommittees of every size across 20 systems, halving the committee to $m = 4$ retains 0.970 of Factor A on average. But the worst individual four-member subcommittee retains only 0.732, and at $m = 2$, 0.548. Committee size buys *insurance*, not accuracy: going from four heads to eight moves the mean by three points and the floor by twenty-seven.

7.2 INDEPENDENTLY TRAINED COMMITTEES

Eight separately trained models, matched to shared-trunk counterparts on training data and evaluation frames, confirm the global-statistic recommendation—so it is a property of committees, not of the multi-head architecture. They are also far better calibrated in magnitude: overconfidence $2.28\times$ against $25.77\times$ for the matched shared-trunk committee. The *acceleration*, by contrast, does not transfer at all, and we state this as a derivation rather than a measurement: with no shared trunk,




Figure 4: The global statistic is extensive, and normalisation restores parity rather than superiority. (a) Median atom count of the acquired set on the MPtraj pool. The global statistic and the two quantities that inherit its extensivity acquire far larger structures than the pool median; the intensive statistics and the free signal do not. (b) Held-out AUROC of $S_\alpha = (\sum_d v_d)/(3N)^\alpha$ against the exponent α , with the design-selected α^* starred. The curve peaks near max-component’s level, not above it.

obtaining Fw requires all M columns of F , so no sketch can be cheaper than the exact quantity. Of five pre-registered predictions in this experiment, 3 were confirmed.

7.3 CONDENSED PHASE AND DISTRIBUTION SHIFT

Bulk liquid water ($N = 192$, periodic, $D = 576$) and a same-temperature MD-versus-PIMD pair give a distribution shift in which only quantum nuclear effects differ. One caveat must survive contact with the reader: the shift regime’s striking Factor A is substantially a symptom of *this model family failing badly there*. A matched independently-trained committee is $2.7\times$ more accurate on the shifted data and its Factor A collapses, while a free signal retains most of the apparent skill. Shift-regime claims must therefore be quoted beside the free-signal baseline and the accuracy gap.

8 FOUNDATION MODELS, REAL POOLS, AND THE BASELINE THAT MUST BE BEATEN

The setting the method is actually pitched at is a foundation model with committee readouts on a frozen trunk (Batatia et al., 2023). We scan 136,923 periodic crystals drawn from MPtraj (Deng et al., 2023), spanning 89 elements and 1–444 atoms, with a MACE-MP (Batatia et al., 2022) trunk fine-tuned to carry eight disjoint committee heads. Five predictions were pre-registered; two were confirmed.

8.1 FACTOR A TRANSFERS

Exact committee disagreement reaches AUROC 0.798 here against a random null of 0.499. Committee force uncertainty is not an artefact of small organic molecules.

8.2 EXTENSIVITY, AND WHY THE GLOBAL STATISTIC REVERSES

The global statistic sums v_d over $3N$ coordinates, so it grows with system size whether or not the model is uncertain. Acquiring by it returns structures whose median atom count is $2.50\times$ that of the max-component selection, and the ordering of Section 4 reverses on all four error scores (Figure 4a, Table 1). Every system in Section 4 has fixed N , where an extensive and an intensive statistic differ by one constant and the pathology *cannot* appear: the recommendation was never wrong where it was measured, it was measured only where it could not fail.

Table 1: On a real materials pool a free signal outranks every committee statistic. AUROC on 136,923 MPtrajectories under a frozen foundation-model trunk. The last two columns are post-hoc diagnostics: AUROC recomputed within atom-count quartiles and within force-magnitude quintiles, so neither the extensivity confound nor force magnitude itself can explain the ordering. $\|\tilde{f}\|$ is bold because it is the signal that costs nothing.

signal	AUROC by error score				stratified (post hoc)	
	$e_{\max}$	e_{rmse}	$e_{\max c}$	e_{q95}	by size	by force
exact max-atom	0.836	0.819	0.824	0.806	–	0.649
exact max-component	0.821	0.808	0.822	0.805	0.798	0.656
exact global	0.798	0.769	0.793	0.759	0.794	0.632
leading-only, 5 lanes	0.791	0.763	0.786	0.753	0.779	0.626
$\ \tilde{f}\ $ (<i>free</i>)	0.943	0.933	0.942	0.930	0.933	0.710
energy std (<i>free</i>)	0.481	0.447	0.483	0.450	0.394	0.511
random null	0.499	0.498	0.498	0.498	–	0.499

Table 2: Pre-registered systems predictions, scored per device on matched pairs. Water and 3BPA measured in one exclusive job per device at identical settings. v2 is the frozen pre-registered protocol; v3 is a post-hoc repair that derives its reference from the matched 3BPA run rather than a fixed constant, and it changes the score only on Blackwell. P1 and P2 — the predictions carrying the speedup claim — pass everywhere.

device	verdict		pre-registered predictions				speedup
	v2	v3	P1	P2	P3	P4	water, $B=16$
H200	4/4	4/4	✓	✓	✓	✓	1.852×
A100-SXM4-80GB	3/4	3/4	✓	✓	×	✓	1.880×
RTX PRO 6000 Blackwell	2/4	3/4	✓	✓	×	×	1.894×

The repair is a size normalisation, $S_\alpha = (\sum_d v_d)/(3N)^\alpha$, with α chosen on a design split and every reported number taken from the disjoint remainder. The selected $\alpha = 0.75$ raises held-out AUROC from 0.797 to 0.821 (Figure 4b), and at fixed N it changes nothing exactly—so it can be adopted globally at no cost. We report the outcome precisely: across four error scores the normalised statistic achieves 2 significant wins, one tie, and 1 significant loss against max-component. This restores *parity*, not superiority, and α is pool- and target-dependent rather than universal.

8.3 THE FREE SIGNAL WINS HERE, AND IS NOT COMPLEMENTED

The max mean-force norm reaches AUROC 0.943 against the committee’s 0.798—a signal that is already computed in the forward pass outranking eight backward passes (Table 1). The advantage survives stratification by atom count and, within force-magnitude quintiles, the committee still beats chance, so it does carry information beyond force magnitude; a free proxy is simply at least as good.

Because the free signal is free, the decision a practitioner faces is not “committee or free signal” but “free signal, or free signal *plus* eight backward passes”. Nested models fitted on a design split and scored on held-out data give AUROC 0.9425 for the free signal alone and 0.9427 with the committee added—a marginal of 0.0002, not significant. On fixed-composition systems the same comparison gives significant gains on 8 of 10 systems, spanning 0.022–0.243. The committee is redundant *here*, not in general, and the discriminating variable is how close the free signal already is to saturation.

8.4 THE SELECTION-OVERLAP BOUND IS UNAVAILABLE

One can hope to bound the downstream active-learning difference without retraining, by showing the gate acquires nearly the same set as the exact statistic. It does not: Jaccard overlap is 0.794 at a realistic acquisition rate, short of the pre-registered threshold. Overlap statistics must always be reported with their acquisition rate, since any two rankings agree when a large fraction of the pool is taken.




Figure 5: Lane reduction pays once the backward pass is compute-bound, and the batch size at which that happens rises with machine speed. Total speedup against batch size for (a) water and (b) 3BPA, measured with both systems in a single exclusive job per device. Water converges to $1.88\times$ – $1.89\times$ by $B = 16$ on all three architectures; 3BPA stays near $1.10\times$ except on Blackwell. The H200 is the fastest machine and is still launch-bound at $B = 1$.

9 WALL CLOCK

Lane reduction pays only once the backward pass is compute-bound rather than launch-bound, and the controlling variable is the total edge count in the batch, not the number of structures. Measuring both systems in a single exclusive job at identical settings, water reaches $1.53\times$ at $B = 1$ rising to $1.88\times$ at $B = 16$ on an A100, $1.85\times$ on an H200 and $1.89\times$ on an RTX PRO 6000 Blackwell—three GPU generations agreeing to 2.3%—against $1.10\times$ for 3BPA in the same job (Figure 5, Table 2). Because the threshold is a property of the machine as well as the system, the batch size at which lane reduction begins to pay *increases* with how fast the GPU is.

Acceleration is deliberately not the headline. The count of avoided exact evaluations is hardware-independent and survives every timing outcome.

10 RECOMMENDATIONS AND LIMITATIONS

1. **Score committee uncertainty against the free mean-force norm**, not against a random null. On some system classes it wins by a wide margin; on others it adds nothing measurable.
2. **Use the global statistic at fixed system size**; normalise by $(3N)^\alpha$ for variable size and refit α per pool and per target.
3. **Prefer a deterministic leading subspace to a randomised sketch**, at equal budget.
4. **Committee size buys insurance**. To fix magnitude calibration, change the construction, not M .

Limitations. One architecture family; $M = 8$ with subcommittee interpolation and no trained $M = 16$ control. No retraining was performed, so no claim is made about downstream model quality—only about which structures are selected. The selection-overlap bound is unavailable, so that gap cannot be closed cheaply. α is pool-dependent. Finally, several results here correct earlier statements of our own; the record of those corrections is retained in the repository rather than silently overwritten.

REPRODUCIBILITY STATEMENT / DATA AND SOFTWARE AVAILABILITY

All statistical results regenerate on CPU from repository-tracked artifacts: `tools/reproduce.py` -tier offline re-derives every headline quantity in roughly four minutes with no model execution, no CUDA and no download, and checks each against its committed value. `tools/audit.py` verifies the frozen prior-work tree, the sha256 binding of every pre-registered protocol to a registration record written before that experiment produced any artifact, the provenance of every number in this manuscript, the split manifests, and the test-access ledger. `tools/report_go_no_go.py` emits every scored verdict as data, alongside the confidence stated for it in advance.

USE OF LARGE LANGUAGE MODELS

LLM-based coding agents were used substantially throughout this work: for research ideation, for implementing the experiment scripts and analysis code, for the adversarial review described below, and for drafting and revising this manuscript. The author directed the work, chose which questions to ask, and is responsible for its contents.

Because that exposure is broad, the verification is mechanical rather than a matter of inspection. Every number in this paper is generated from a committed record by `tools/paper_numbers.py` and re-checked by `tools/audit.py`; none is transcribed, so an agent cannot introduce a figure by typing one. Each experiment’s protocol was registered, and hash-bound to a record written before that experiment produced any artifact, so a hypothesis cannot be adjusted after seeing the outcome. Split roles are content-hashed and recompute from recorded seeds.

That machinery exists because agent-written prose did fail in exactly the way one would predict. An adversarial review, itself agent-driven and documented in `results/REVIEW_FINDINGS.md`, confirmed 22 real defects, two of them critical: a benchmark that silently measured seven lanes where it reported eight, and a filename collision that destroyed a result file. A separate cross-document audit found roughly 70 quantitative claims in the supporting documents that a record contradicts, nearly all of them stale copies left behind by corrections that were made in one place and not propagated. The macro-derived numbers — the ones in this manuscript — were not among the failures, which is the entire argument for deriving them.

REFERENCES

- Ilyes Batatia, Dávid Péter Kovács, Gregor N. C. Simm, Christoph Ortner, and Gábor Csányi. MACE: Higher order equivariant message passing neural networks for fast and accurate force fields. In *Advances in Neural Information Processing Systems 35*, pp. 11423–11436, 2022. doi: 10.52202/068431-0830.
- Ilyes Batatia, Philipp Benner, Yuan Chiang, et al. A foundation model for atomistic materials chemistry. arXiv:2401.00096 [physics.chem-ph], 2023. doi:10.48550/arXiv.2401.00096.
- Hubert Beck, Pavol Simko, Lars L. Schaaf, Ondrej Marsalek, and Christoph Schran. Multi-head committees enable direct uncertainty prediction for atomistic foundation models. *The Journal of Chemical Physics*, 163(23):234103, 2025. doi: 10.1063/5.0302097. arXiv:2508.09907.
- Anders S. Christensen and O. Anatole von Lilienfeld. On the role of gradients for machine learning of molecular energies and forces. *Machine Learning: Science and Technology*, 1(4):045018, 2020. doi: 10.1088/2632-2153/abba6f.
- Bowen Deng, Peichen Zhong, KyuJung Jun, Janosh Riebesell, Kevin Han, Christopher J. Bartel, and Gerbrand Ceder. CHGNet as a pretrained universal neural network potential for charge-informed atomistic modelling. *Nature Machine Intelligence*, 5(9):1031–1041, 2023. doi: 10.1038/s42256-023-00716-3.
- Dávid Péter Kovács, Cas van der Oord, Jiri Kucera, Alice E. A. Allen, Daniel J. Cole, Christoph Ortner, and Gábor Csányi. Linear atomic cluster expansion force fields for organic molecules: Beyond RMSE. *Journal of Chemical Theory and Computation*, 17(12):7696–7711, 2021. doi: 10.1021/acs.jctc.1c00647.

- Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics*, 17(3), 1989. doi: 10.1214/aos/1176347265.
- Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523): 1094–1111, 2018. doi: 10.1080/01621459.2017.1307116.
- Raphael A. Meyer, Cameron Musco, Christopher Musco, and David P. Woodruff. Hutch++: Optimal stochastic trace estimation. In *Symposium on Simplicity in Algorithms (SOSA)*, pp. 142–155. Society for Industrial and Applied Mathematics, 2021. doi: 10.1137/1.9781611976496.16.
- H. S. Seung, M. Opper, and H. Sompolinsky. Query by committee. In *Proceedings of the Fifth Annual Workshop on Computational Learning Theory*, pp. 287–294, 1992. doi: 10.1145/130385.130417.
- Vladimir Vovk, Alexander Gammernan, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005. doi: 10.1007/b106715.